

# Exam notes

## Functional dependencies

**Trivial FD:**  $X \rightarrow Y$  is trivial if:  $Y \subset X$ .

**Non-trivial FD:** see above.

**Superkey:** Superset of a candidate key.

**Candidate key:** Minimal set of a key.

**Primary key:** Candidate key chosen by designer.

**Prime attributes:** Attributes contained in candidate key.

**Closure of FDs:** Set of all dependencies that include  $F$  and all those that can be inferred from  $F$ . Denoted as  $F^+$ .

**Closure of set of attributes:** Maximum set of attributes that can be inferred from the provided set of attributes.

## Normalization

Database design principle to

1. *reduce redundancy*
2. *avoid complexities*
3. *organize data in a consistent way.*
4. *eliminate duplicates*  
but loses out on

- query efficiency

**1NF:** No multivalued attributes, i.e. only simple values.

**2NF:** Parts of the candidate key must not be functionally dependent on non-prime attributes.

- **Example:**  $\{\mathbf{A}, \mathbf{B}\} : \overbrace{\mathbf{A} \rightarrow C, D, E}^{\text{illegal}}$   
 $\Rightarrow$  No arrows from part of candidate key!

**Decompose:** Make new table containing LHS and RHS. Remove RHS from origin table.

**3NF:** Non-prime attributes must not be functionally determined by other non-prime attributes.

- Equivalent to no transitive dependencies:  $A \rightarrow \overbrace{B \rightarrow C}^{\text{illegal}}$   
 $\Rightarrow$  No arrows starting from non-prime attributes!

**Decompose:** Make new table containing LHS and RHS. Remove RHS from origin table.

**BCNF:** No dependencies from non-key attributes to key attributes.

- Example:  $\{A, B\} \rightarrow \overbrace{C \rightarrow A}^{\text{illegal}}$   
 $\Rightarrow$  No arrows pointing towards prime attributes!

**Decompose:** Remove one of the prime-attributes!

**NB:** BCNF is lossless, but it is not dependency preserving!

**Losslessness:** Decomposition is lossless if we can recover initial table by performing multiple joins.

**Dependency preservation:** We do not lose dependencies. True if all dependencies can be inferred from the current set.

## Triggers

- Can not modify the database schema; solely DML (*Data Manipulation Language*) for data-level operations.

Invoked *automatically*.

Defined in terms of the event *that invokes it*, and the *action it performs*.

May operate either BEFORE or AFTER the execution of the event that invokes it.

## Storage

Main memory vs. disk:

- Data access from disk is typically 2 orders of magnitude slower.
- Data in main memory is **volatile**, while **non-volatile** for disk.
- Disk is solely mechanical moving parts, while main memory is completely electronic.

Each time we read from disk, we retrieve a *block* of records.

Size of these *blocks* is fixed, but depends on the OS.

## Indexes

**Index:** data structure that facilitates quicker access to a data.

- *can be* used for both main memory and disk!
- *stored in a data file*.

**Primary index:** Indexes on the primary key (ordered on key).

- --> **Dense index**: Has exactly one index entry for each search key value.
- --> **Sparse index**: Has fewer index entries than search key values.

**Clustering index**: Indexes on a non-key field (ordered on field).

- --> One index entry of each (distinct!) value of the field.
- --> Each index points to first data block of records for search key.

**Secondary index**: Not ordered on the index's search key (purely a mess!)

**Multi-level index**: Structure of index on index until all entries of the top-level structure fit into 1 disk block.

- --> Pin top-level index in main memory (RAM).

## **$B^+$ -trees**

- *multi-leveled* indexing structure
- tailored for disk-based data organization: *aligns with disk block sizes, so very efficient for disk storage and access.*
- grows horizontally by splitting the root.

| ↓ Operation | Average     | Worst case  |
|-------------|-------------|-------------|
| Search      | $O(\log n)$ | $O(\log n)$ |
| Insert      | $O(\log n)$ | $O(\log n)$ |
| Delete      | $O(\log n)$ | $O(\log n)$ |

**Balanced**: all paths from root to a leaf have the same length.

- --> guarantees good search performance!

## **Recovery strategies**

**Checkpointing**: process where the system periodically writes all modified (*dirty*) pages from memory to disk.

- improves **recovery** efficiency
- ensures a **consistent** state can be restored after unexpected crash.

**Shadow-paging**: copy-on-write technique for avoiding in-place updates of pages

- when a page is modified (*dirty*), the system writes changes to a **new (shadow) page** instead of overwriting the old.

- upon commit the page table pointer is switched to the new page (**atomic!**).  
≈ no-undo / no-redo

⇒ possible to use both techniques, but often redundant because both handle recovery well.

**Undo:** rollback changes of *uncommitted* transactions

**Redo:** reapply changes of *committed* transactions after a crash.

**Steal:** *uncommitted* dirty pages can be written, so you need to **undo** them.

- e.g. to save space in main memory you flush the dirty pages more often.

**Force:** *committed* pages are immediately written (forced!) to disk, so you do not need to repeat them (**no-redo**).

- generally not favorable because it leads to a ton of continuously costly I/Os

|          | Steal          | No-steal          |
|----------|----------------|-------------------|
| Force    | Undo / no-redo | No-undo / no-redo |
| No-force | Undo / redo    | No-undo / redo    |

**Undo/redo:**

1. Undo all transactions that has a log entry of "start" but no "commit".
2. Redo all transactions that has a long entry of "start" and "commit".

## Misc