



Advanced SQL

Databases



Ira Assent

ira@cs.au.dk

Data-Intensive Systems group, Department of Computer Science, Aarhus University, DK

Intended learning outcomes

- ▶ Be able to
 - ▶ Write simple SQL DML queries
 - ▶ Make use of advanced constructs in DML queries

Recap: SQL DDL

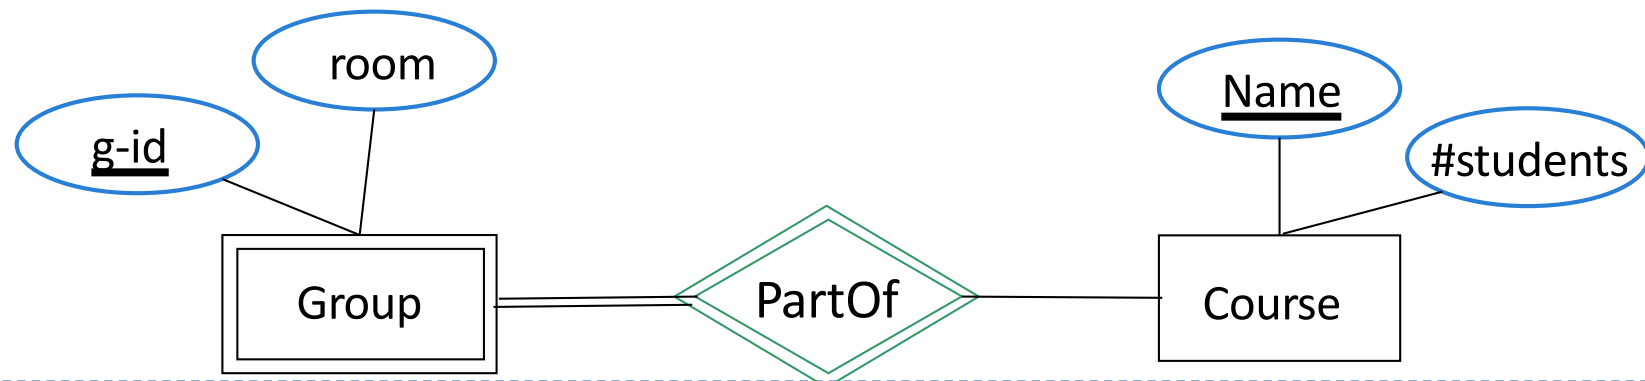
- ▶ **DDL: define database schema in SQL**
- ▶ `CREATE DATABASE Uni;`
- ▶ `CREATE TABLE Student (id INT PRIMARY KEY, name VARCHAR(15), office VARCHAR(10) REFERENCES Rooms(rooms));`
- ▶ **Data types for numeric, string, date etc**
- ▶ **If you want quality (and you do), define constraints!**
 - ▶ **Keys: uniquely identify tuples**
 - ▶ Primary key: **within a table**
 - ▶ Foreign key: **across tables**
- ▶ **Other properties:** NOT NULL, UNIQUE, DEFAULT, triggered actions CASCADE, SET NULL, SET DEFAULT
- ▶ **Make use of documentation for definitions and examples**



<https://dev.mysql.com/doc/refman/8.0/en/create-table.html>

Which SQL DDL for Group?

- A. `CREATE TABLE GROUP(gid INT PRIMARY KEY, room VARCHAR(10));`
- B. `CREATE TABLE GROUP(gid INT PRIMARY KEY, Name VARCHAR(15) PRIMARY KEY, room VARCHAR(10));`
- C. `CREATE TABLE GROUP(gid INT, Name VARCHAR(15), room VARCHAR(10), PRIMARY KEY(gid,Name));`
- D. `CREATE TABLE GROUP(gid INT PRIMARY KEY, Name VARCHAR(15), room VARCHAR(10), FOREIGN KEY(Name) REFERENCES Course(Name));`
- E. `CREATE TABLE GROUP(gid INT, Name VARCHAR(15), room VARCHAR(10), PRIMARY KEY (gid, Name), FOREIGN KEY(Name) REFERENCES Course(Name));`



SQL DML: SELECT-FROM-WHERE

- ▶ The basic form of an SQL query

```
SELECT desired attributes  
FROM one or more tables  
WHERE condition about the involved rows
```

topic

DB

DB

Meetings

| meetid | date | owner | topic |
|--------|------------|-------|-------|
| 34716 | 2023-08-28 | ira | DB |
| 34717 | 2024-01-22 | ira | DB |
| 42835 | 2023-08-18 | aas | Prog |

- ▶ Which meetings
("topic") has ira arranged?

```
SELECT topic  
FROM Meetings  
WHERE owner = 'ira';
```

Loop Semantics for Single Table

- ▶ Loop through all rows in the table
- ▶ Check if the condition is true
- ▶ Project the rows onto the desired attributes

Note that duplicates are kept!

(Funny) terminology

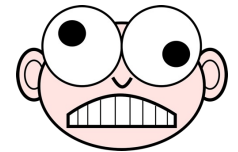
- ▶ The basic form of an SQL query

```
SELECT desired attributes  
FROM one or more tables  
WHERE condition about the involved rows
```

Projection: projecting to attributes

Join: combining tables

Selection condition:
selecting rows



- ▶ Which meetings
("topic") has ira arranged?

```
SELECT topic  
FROM Meetings  
WHERE owner = 'ira';
```

Meetings

| meetid | date | owner | topic |
|--------|------------|-------|-------|
| 34716 | 2023-08-28 | ira | DB |
| 34717 | 2024-01-22 | ira | DB |
| 42835 | 2023-08-18 | aas | Prog |

topic

DB

DB

Combining conditions in WHERE

- ▶ AND, OR, NOT, =, <>, <, >, <=, >=, ...

```
SELECT owner, topic  
FROM Meetings  
WHERE owner = 'ira' OR topic = 'DB' ;
```

Meetings

| meetid | date | owner | topic |
|--------|------------|-------|-------|
| 34716 | 2023-08-28 | ira | DB |
| 34717 | 2024-01-22 | ira | DB |
| 42835 | 2023-08-18 | aas | Prog |

| owner | topic |
|-------|-------|
| ira | DB |
| Ira | DB |

Renaming in SELECT

- ▶ The selected attributes can be given new names

```
SELECT name AS navn, office AS kontor  
FROM People  
WHERE office = 'Ny-357';
```

| userid | name | office |
|--------|-----------------|--------|
| ira | Ira Assent | Ny-357 |
| aas | Annika Schmidt | NULL |
| jan | Jan Christensen | Ho-017 |

| navn | kontor |
|------------|--------|
| Ira Assent | vip |

Renaming in FROM

► Example

```
SELECT name, office
      FROM People AS Folk
 WHERE Folk.office = 'Ny-357';
```

| userid | name | office |
|--------|-----------------|--------|
| ira | Ira Assent | Ny-357 |
| aas | Annika Schmidt | NULL |
| jan | Jan Christensen | Ho-017 |

| name | office |
|------------|--------|
| Ira Assent | vip |



General renaming options

- ▶ Specifying a new name right after the original one using AS called an **alias** or **tuple variable**
 - ▶ Declare alternative relation name E instead of Employee
 - ▶ `SELECT Fname FROM EMPLOYEE AS E;`
 - ▶ Has the same attributes and entries
 - ▶ `SELECT Fn FROM EMPLOYEE AS E(Fn, Mi, Ln, Ssn, Bd, Addr, Sex, Sal, Sssn, Dno);`
 - ▶ Also changes the attribute names

EMPLOYEE

| Fname | Minit | Lname | <u>Ssn</u> | Bdate | Address | Sex | Salary | Super_ssn | Dno |
|----------|-------|---------|------------|------------|--------------------------|-----|--------|-----------|-----|
| John | B | Smith | 123456789 | 1965-01-09 | 731 Fondren, Houston, TX | M | 30000 | 333445555 | 5 |
| Franklin | T | Wong | 333445555 | 1955-12-08 | 638 Voss, Houston, TX | M | 40000 | 888665555 | 5 |
| Alicia | J | Zelaya | 999887777 | 1968-01-19 | 3321 Castle, Spring, TX | F | 25000 | 987654321 | 4 |
| Jennifer | S | Wallace | 987654321 | 1941-06-20 | 291 Berry, Bellaire, TX | F | 43000 | 888665555 | 4 |
| Ramesh | K | Narayan | 666884444 | 1962-09-15 | 975 Fire Oak, Humble, TX | M | 38000 | 333445555 | 5 |
| Joyce | A | English | 453453453 | 1972-07-31 | 5631 Rice, Houston, TX | F | 25000 | 333445555 | 5 |
| Ahmad | V | Jabbar | 987987987 | 1969-03-29 | 980 Dallas, Houston, TX | M | 25000 | 987654321 | 4 |
| James | E | Borg | 888665555 | 1937-11-10 | 450 Stone, Houston, TX | M | 55000 | NULL | 1 |

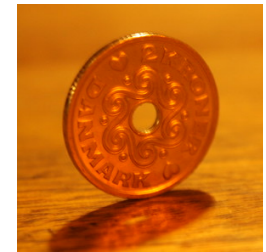
Math in SELECT

- ▶ The attributes may have computed values

```
SELECT customer AS kunde, date AS dato, price*7.44 AS pris  
FROM Purchase  
WHERE customer = 'ira';
```

Allows us to convert the price from EUR to DKK

| id | date | price | customer | product |
|----|------------|-------|----------|---------|
| 16 | 2024-02-03 | 2.99 | ira | Sencha |
| 17 | 2024-01-22 | 1.59 | ira | Müsli |
| 5 | 2023-12-18 | 6.49 | aas | Arabica |



| kunde | dato | pris |
|-------|------------|---------|
| ira | 2024-02-03 | 22.2456 |
| ira | 2024-01-22 | 11.8296 |

Arithmetic Operators

- ▶ Standard arithmetic operators:
 - ▶ Addition (+), subtraction (−), multiplication (*), and division (/)
 - ▶ Convenient both in SELECT clause and in WHERE clause

```
SELECT customer AS PricelessCustomer
```

```
FROM Purchase
```

```
WHERE fee + price > 7;
```

| id | fee | price | customer |
|----|------|-------|----------|
| 16 | 0.15 | 2.99 | ira |
| 17 | 2.5 | 1.59 | ira |
| 5 | 1.75 | 6.49 | aas |

| customer |
|----------|
| aas |

- ▶ BETWEEN comparison operator
 - ▶ Returns values that are greater than or equal to first number and smaller than or equal to second number

```
SELECT price AS Sweetspot
```

```
FROM Purchase
```

```
WHERE price BETWEEN 2.99 AND 4.99;
```

| price |
|-------|
| 2.99 |

Multiple Relations / Joins

- ▶ Who has booked meetings on January 22, 2024?

```
SELECT name
FROM People, Meetings
WHERE date = '2024-01-22'
AND owner = userid;
```

- ▶ Relations are **joined**

- ▶ **Join condition** `owner=userid` states how attributes from joined tables should match

General Loop Semantics

- ▶ Loop through all rows in all tables
- ▶ For each combination
 - ▶ check if the condition is true
 - ▶ project the rows onto the desired attributes
- ▶ Note that duplicates are still kept

| Meetings | | | | People | | | |
|----------|------------|-------|------------|--------|-----------------|-------|--------|
| meetid | date | owner | topic | userid | name | group | office |
| 34716 | 2023-08-28 | ira | dDB | ira | Ira Assent | vip | Ny-357 |
| 34717 | 2024-01-22 | ira | dDB | aas | Annika Schmidt | phd | NULL |
| 42835 | 2023-08-18 | aas | TA meeting | jan | Jan Christensen | tap | Ho-017 |

| name |
|------------|
| Ira Assent |

Example Company database from textbook

EMPLOYEE

| Fname | Minit | Lname | <u>Ssn</u> | Bdate | Address | Sex | Salary | Super_ssn | Dno |
|----------|-------|---------|------------|------------|--------------------------|-----|--------|-----------|-----|
| John | B | Smith | 123456789 | 1965-01-09 | 731 Fondren, Houston, TX | M | 30000 | 333445555 | 5 |
| Franklin | T | Wong | 333445555 | 1955-12-08 | 638 Voss, Houston, TX | M | 40000 | 888665555 | 5 |
| Alicia | J | Zelaya | 999887777 | 1968-01-19 | 3321 Castle, Spring, TX | F | 25000 | 987654321 | 4 |
| Jennifer | S | Wallace | 987654321 | 1941-06-20 | 291 Berry, Bellaire, TX | F | 43000 | 888665555 | 4 |
| Ramesh | K | Narayan | 666884444 | 1962-09-15 | 975 Fire Oak, Humble, TX | M | 38000 | 333445555 | 5 |
| Joyce | A | English | 453453453 | 1972-07-31 | 5631 Rice, Houston, TX | F | 25000 | 333445555 | 5 |
| Ahmad | V | Jabbar | 987987987 | 1969-03-29 | 980 Dallas, Houston, TX | M | 25000 | 987654321 | 4 |
| James | E | Borg | 888665555 | 1937-11-10 | 450 Stone, Houston, TX | M | 55000 | NULL | 1 |

DEPARTMENT

| Dname | <u>Dnumber</u> | Mgr_ssn | Mgr_start_date |
|----------------|----------------|-----------|----------------|
| Research | 5 | 333445555 | 1988-05-22 |
| Administration | 4 | 987654321 | 1995-01-01 |
| Headquarters | 1 | 888665555 | 1981-06-19 |

DEPT_LOCATIONS

| <u>Dnumber</u> | <u>Dlocation</u> |
|----------------|------------------|
| 1 | Houston |
| 4 | Stafford |
| 5 | Bellaire |
| 5 | Sugarland |
| 5 | Houston |

Simple queries on the example

Query 0. Retrieve the birth date and address of the employee(s) whose name is 'John B. Smith'.

Q0: **SELECT** Bdate, Address
 FROM EMPLOYEE
 WHERE Fname='John' AND Minit='B' AND Lname='Smith';

| <u>Bdate</u> | <u>Address</u> |
|--------------|--------------------------|
| 1965-01-09 | 731 Fondren, Houston, TX |

Query 1. Retrieve the name and address of all employees who work for the 'Research' department.

Q1: **SELECT** Fname, Lname, Address
 FROM EMPLOYEE, DEPARTMENT
 WHERE Dname='Research' AND Dnumber=Dno;

| <u>Fname</u> | <u>Lname</u> | <u>Address</u> |
|--------------|--------------|--------------------------|
| John | Smith | 731 Fondren, Houston, TX |
| Franklin | Wong | 638 Voss, Houston, TX |
| Ramesh | Narayan | 975 Fire Oak, Humble, TX |
| Joyce | English | 5631 Rice, Houston, TX |

EMPLOYEE

| <u>Fname</u> | <u>Minit</u> | <u>Lname</u> | <u>Ssn</u> | <u>Bdate</u> | <u>Address</u> | <u>Sex</u> | <u>Salary</u> | <u>Super_ssn</u> | |
|--------------|--------------|--------------|------------|--------------|--------------------------|------------|---------------|------------------|---|
| John | B | Smith | 123456789 | 1965-01-09 | 731 Fondren, Houston, TX | M | 30000 | 333445555 | 5 |
| Franklin | T | Wong | 333445555 | 1955-12-08 | 638 Voss, Houston, TX | M | 40000 | 888665555 | 5 |
| Alicia | J | Zelaya | 999887777 | 1968-01-19 | 3321 Castle, Spring, TX | F | 25000 | 987654321 | 4 |
| Jennifer | S | Wallace | 987654321 | 1941-06-20 | 291 Berry, Bellaire, TX | F | 43000 | 888665555 | 4 |
| Ramesh | K | Narayan | 666884444 | 1962-09-15 | 975 Fire Oak, Humble, TX | M | 38000 | 333445555 | 5 |
| Joyce | A | English | 453453453 | 1972-07-31 | 5631 Rice, Houston, TX | F | 25000 | 333445555 | 5 |
| Ahmad | V | Jabbar | 987987987 | 1969-03-29 | 980 Dallas, Houston, TX | M | 25000 | 987654321 | 4 |
| James | E | Borg | 888665555 | 1937-11-10 | 450 Stone, Houston, TX | M | 55000 | NULL | 1 |

DEPARTMENT

| <u>Dname</u> | <u>Dnumber</u> | <u>Mgr_ssn</u> | <u>Mgr_start_date</u> |
|----------------|----------------|----------------|-----------------------|
| Research | 5 | 333445555 | 1988-05-22 |
| Administration | 4 | 987654321 | 1995-01-01 |
| Headquarters | 1 | 888665555 | 1981-06-19 |

DEPT_LOCATIONS

| <u>Dnumber</u> | <u>Dlocation</u> |
|----------------|------------------|
| 1 | Houston |
| 4 | Stafford |
| 5 | Bellaire |

Joining more tables

Query 2. For every project located in 'Stafford', list the project number, the controlling department number, and the department manager's last name, address, and birth date.

Q2: **SELECT** Pnumber, Dnum, Lname, Address, Bdate
 FROM PROJECT, DEPARTMENT, EMPLOYEE
 WHERE Dnum=Dnumber **AND** Mgr_ssn=Ssn **AND**
 Plocation='Stafford';

| <u>Pnumber</u> | <u>Dnum</u> | <u>Lname</u> | <u>Address</u> | <u>Bdate</u> |
|----------------|-------------|--------------|-------------------------|--------------|
| 10 | 4 | Wallace | 291 Berry, Bellaire, TX | 1941-06-20 |
| 30 | 4 | Wallace | 291 Berry, Bellaire, TX | 1941-06-20 |

EMPLOYEE

| Fname | Minit | Lname | <u>Ssn</u> | Bdate | Address | Sex | Salary | Super_ssn | Dno |
|----------|-------|---------|------------|------------|--------------------------|-----|--------|-----------|-----|
| John | B | Smith | 123456789 | 1965-01-09 | 731 Fondren, Houston, TX | M | 30000 | 333445555 | 5 |
| Franklin | T | Wong | 333445555 | 1955-12-08 | 638 Voss, Houston, TX | M | 40000 | 888665555 | 5 |
| Alicia | J | Zelaya | 999887777 | 1968-01-19 | 3321 Castle, Spring, TX | F | 25000 | 987654321 | 4 |
| Jennifer | S | Wallace | 987654321 | 1941-06-20 | 291 Berry, Bellaire, TX | F | 43000 | 888665555 | 4 |
| Ramesh | K | Narayan | 666884444 | 1962-09-15 | 975 Fire Oak, Humble, TX | M | 38000 | 333445555 | 5 |
| Joyce | A | English | 453453453 | 1972-07-31 | 5631 Rice, Houston, TX | F | 25000 | 333445555 | 5 |
| Ahmad | V | Jabbar | 987987987 | 1969-03-29 | 980 Dallas, Houston, TX | M | 25000 | 987654321 | 4 |
| James | E | Borg | 888665555 | 1937-11-10 | 450 Stone, Houston, TX | M | 55000 | NULL | 1 |

PROJECT

| Pname | <u>Pnumber</u> | Plocation | Dnum |
|-----------------|----------------|-----------|------|
| ProductX | 1 | Bellaire | 5 |
| ProductY | 2 | Sugarland | 5 |
| ProductZ | 3 | Houston | 5 |
| Computerization | 10 | Stafford | 4 |
| Reorganization | 20 | Houston | 1 |
| Newbenefits | 30 | Stafford | 4 |

DEPARTMENT

| Dname | <u>Dnumber</u> | Mgr_ssn | Mgr_start_date |
|----------------|----------------|-----------|----------------|
| Research | 5 | 333445555 | 1988-05-22 |
| Administration | 4 | 987654321 | 1995-01-01 |
| Headquarters | 1 | 888665555 | 1981-06-19 |

DEPT_LOCATIONS

| <u>Dnumber</u> | <u>Dlocation</u> |
|----------------|------------------|
| 1 | Houston |
| 4 | Stafford |
| 5 | Bellaire |

Which of these is correct?

1. `SELECT name FROM Meetings, Rooms WHERE capacity = '6' AND room = office;`
2. `SELECT Meetings, Rooms WHERE capacity = '6' AND room = office;`
3. `SELECT name FROM Meetings AND Rooms WHERE capacity = '6' AND room = office;`
4. `SELECT name FROM Meetings WHERE capacity = '6' AND room = office;`

| userid | name | group | office |
|--------|-----------------|-------|--------|
| ira | Ira Assent | vip | Ny-357 |
| aas | Annika Schmidt | phd | NULL |
| jan | Jan Christensen | tap | Ho-017 |

Meetings

| room | capacity |
|----------|----------|
| Ny-357 | 6 |
| Ada-333 | 26 |
| StoreAud | 286 |

Rooms

Ambiguous Attribute Names

- ▶ Same name can be used for two (or more) attributes
 - ▶ As long as the attributes are in different relations
 - ▶ In practice very common: e.g. attribute names id, name in Student (id, name, studies, city), Professor (id, name, department, teaches)
 - ▶ If using two same name attributes in the same query, must **qualify** attribute name with relation name to prevent ambiguity
 - ▶ `Relation.Attribute`
 - ▶ Else get error message

```
SELECT  Fname, EMPLOYEE.Name, Address
FROM    EMPLOYEE, DEPARTMENT
WHERE   DEPARTMENT.Name='Research' AND
        DEPARTMENT.Dnumber=EMPLOYEE.Dnumber;
```

Unspecified WHERE Clause

- ▶ Missing WHERE clause
 - ▶ Indicates no condition on tuple selection
 - ▶ Thus, all tuples in relation returned
 - ▶ Use as quick view of your table contents!
- ▶ When querying two or more tables returns the cross product
 - ▶ All possible combinations of tuples from either table
 - ▶ i.e., a lot!!



```
SELECT    Ssn
FROM      EMPLOYEE;

SELECT    Ssn, Dname
FROM      EMPLOYEE, DEPARTMENT;
```

```

SELECT      Ssn
FROM        EMPLOYEE;

SELECT      Ssn, Dname
FROM        EMPLOYEE, DEPARTMENT;

```

EMPLOYEE

| Fname | Minit | Lname | <u>Ssn</u> | Bdate | |
|----------|-------|---------|------------|------------|---|
| John | B | Smith | 123456789 | 1965-01-09 | 7 |
| Franklin | T | Wong | 333445555 | 1955-12-08 | 6 |
| Alicia | J | Zelaya | 999887777 | 1968-01-19 | 3 |
| Jennifer | S | Wallace | 987654321 | 1941-06-20 | 2 |
| Ramesh | K | Narayan | 666884444 | 1962-09-15 | 9 |
| Joyce | A | English | 453453453 | 1972-07-31 | 5 |
| Ahmad | V | Jabbar | 987987987 | 1969-03-29 | 9 |
| James | E | Borg | 888665555 | 1937-11-10 | 4 |

DEPARTMENT

| Dname | <u>Dnumber</u> | Mgr_ssn | Mg |
|----------------|----------------|-----------|----|
| Research | 5 | 333445555 | 19 |
| Administration | 4 | 987654321 | 19 |
| Headquarters | 1 | 888665555 | 19 |

| <u>Ssn</u> | <u>Dname</u> |
|------------|----------------|
| 123456789 | Research |
| 333445555 | Research |
| 999887777 | Research |
| 987654321 | Research |
| 666884444 | Research |
| 453453453 | Research |
| 987987987 | Research |
| 888665555 | Research |
| 123456789 | Administration |
| 333445555 | Administration |
| 999887777 | Administration |
| 987654321 | Administration |
| 666884444 | Administration |
| 453453453 | Administration |
| 987987987 | Administration |
| 888665555 | Administration |
| 123456789 | Headquarters |
| 333445555 | Headquarters |
| 999887777 | Headquarters |
| 987654321 | Headquarters |
| 666884444 | Headquarters |
| 453453453 | Headquarters |
| 987987987 | Headquarters |
| 888665555 | Headquarters |

| <u>Salary</u> |
|---------------|
| 000 |
| 000 |
| 000 |
| 000 |
| 000 |
| 000 |
| 000 |
| 000 |
| 000 |
| 000 |
| 000 |
| 000 |
| 000 |
| 000 |
| 000 |

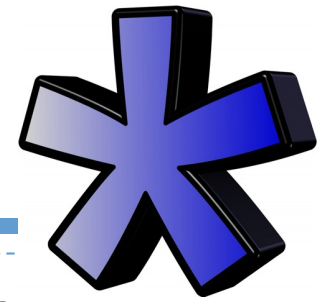
| <u>SSN</u> |
|------------|
| 123456789 |
| 333445555 |
| 999887777 |
| 987654321 |
| 666884444 |
| 453453453 |
| 987987987 |
| 888665555 |

DEPT_LOCATIONS

| <u>location</u> | <u>Dlocation</u> |
|-----------------|------------------|
| 1 | Houston |
| 4 | Stafford |
| 5 | Bellaire |
| 5 | Sugarland |
| 5 | Houston |

How many rows if you do cross product on three tables with 10 tuples each?

Asterisk notation



- ▶ Specify an asterisk * instead of attributes in SELECT clause
 - ▶ Retrieve all the attribute values of the selected tuples
 - ▶ Is the same as listing all attributes of the relation(s)

```
SELECT * FROM EMPLOYEE;
```

same as

```
SELECT (Fname, Minit, Lname, Ssn, Bdate, Address,  
Sex, Salary, Super_ssn, Dno) FROM EMPLOYEE;
```

- ▶ Saves a lot of typing



EMPLOYEE

| Fname | Minit | Lname | <u>Ssn</u> | Bdate | Address | Sex | Salary | Super_ssn | Dno |
|----------|-------|---------|------------|------------|--------------------------|-----|--------|-----------|-----|
| John | B | Smith | 123456789 | 1965-01-09 | 731 Fondren, Houston, TX | M | 30000 | 333445555 | 5 |
| Franklin | T | Wong | 333445555 | 1955-12-08 | 638 Voss, Houston, TX | M | 40000 | 888665555 | 5 |
| Alicia | J | Zelaya | 999887777 | 1968-01-19 | 3321 Castle, Spring, TX | F | 25000 | 987654321 | 4 |
| Jennifer | S | Wallace | 987654321 | 1941-06-20 | 291 Berry, Bellaire, TX | F | 43000 | 888665555 | 4 |

Using the same relation twice (self-join)

Find all pairs of roommates like (Annika, Ira)

Do not produce pairs with self like (Ira, Ira)

Do not produce duplicates like (Ira, Annika) for (Annika, Ira) – use alphabetic order

1. `SELECT name, name FROM People WHERE office = office AND name = name;`
2. `SELECT name, name FROM People WHERE (name, office);`
3. `SELECT p1.name, p2.name FROM People AS p1, People AS p2 WHERE p1.office = p2.office AND p1.name < p2.name;`
4. `SELECT p1.name, p2.name FROM People p1, People p2 WHERE p1.office = p2.office;`

Finding a pattern in a string

We would like to find all events relating to beer that are organized by Frank

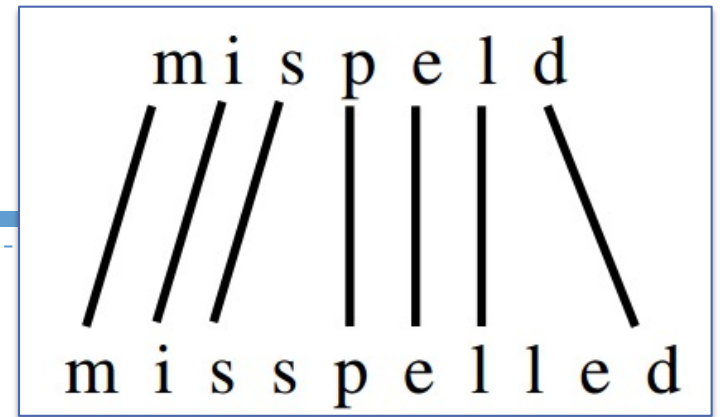
```
SELECT owner, topic
FROM Meetings
WHERE owner = 'fra' AND topic LIKE '%beer%';
```

| owner | topic |
|-------|----------------|
| fra | afternoon beer |
| fra | beer tasting |

Meetings

| meetid | date | owner | topic |
|--------|------------|-------|---------------------|
| 34716 | 2023-08-28 | ira | dDB |
| 34717 | 2024-01-22 | ira | dDB |
| 42835 | 2023-10-18 | aas | TA meeting |
| 43779 | 2024-02-01 | fra | afternoon beer |
| 48333 | 2024-02-08 | fra | beer tasting |
| 50001 | 2024-01-09 | joe | return beer bottles |

Fun with Pattern Matching



- ▶ **LIKE** comparison operator
 - ▶ Used for string **pattern matching**
 - ▶ % replaces an arbitrary number of zero or more characters
 - ▶ underscore **_** replaces a single character

```
SELECT Name FROM Student WHERE Address LIKE '%Aarhus%';
```

- ▶ Matches *Aarhus*, *8200 Aarhus N*, *8000 Aarhus*, etc.

```
SELECT Name FROM Student WHERE Address LIKE 'Aarhus__';
```

- ▶ Matches *Aarhus N*, but **NOT Aarhus!**
- ▶ If you need to find a string with %, precede it with an escape character \

```
SELECT comment FROM Code WHERE comment LIKE '\ \% \ \% \ %%';
```

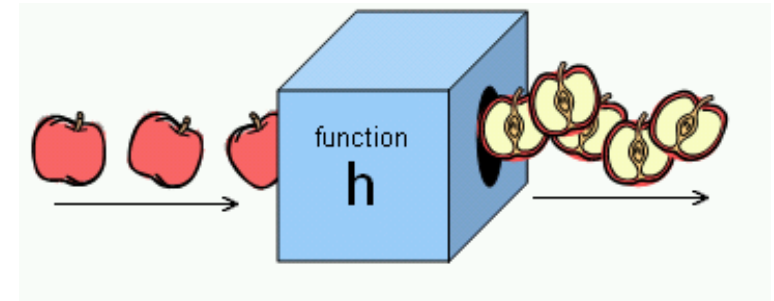
- ▶ Matches *%%%My Comment%%%*

```
SELECT CONCAT(FirstName, ' ', LastName) AS FullName FROM Employees;
```

Concatenates first name, a space, and the last name into a single return value

Scalar functions

- ▶ May use scalar functions
- ▶ Many available
 - ▶ integer and float functions
 - ▶ string functions
 - ▶ calendar functions, ...



```
SELECT CHARACTER_LENGTH(name) as len,  
UPPER(`group`) as `GROUP` FROM People;
```

| userid | name | group | office |
|--------|-----------------|-------|--------|
| ira | Ira Assent | vip | Ny-357 |
| aas | Annika Schmidt | phd | NULL |
| jan | Jan Christensen | tap | Ho-017 |

| len | GROUP |
|-----|-------|
| 9 | VIP |
| 14 | PHD |
| 15 | TAP |

<https://dev.mysql.com/doc/refman/8.0/en/built-in-function-reference.html>

Tables as Sets in SQL

- ▶ SQL does not automatically eliminate duplicate tuples in query results
 - ▶ i.e., uses **multiset semantics**
- ▶ Use the keyword **DISTINCT** in the `SELECT` clause
 - ▶ Turns result into **set**
 - ▶ Each value remains only once in the result
 - ▶ Expensive operation
- ▶ Which studies do our students follow?

```
SELECT studies FROM Student;
```

```
SELECT DISTINCT studies FROM Student;
```

If not distinct, can specify **ALL**

```
SELECT  
FROM
```

```
ALL Salary  
EMPLOYEE;
```

```
SELECT  
FROM
```

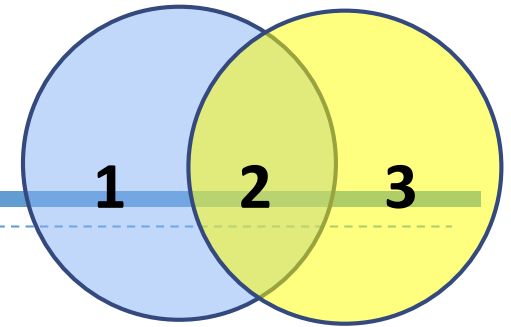
```
DISTINCT Salary  
EMPLOYEE;
```

| id | name | studies |
|----|-------|---------|
| 1 | Joe | CS |
| 2 | Jane | ITP |
| 3 | Alice | CS |
| 4 | Bob | CE |
| 5 | Eve | DS |
| 6 | Bo | CS |
| 7 | Ib | ITP |
| 8 | Liv | ITP |

| studies |
|---------|
| CS |
| ITP |
| CS |
| CE |
| DS |
| CS |
| ITP |
| ITP |

| studies |
|---------|
| CS |
| ITP |
| CE |
| DS |

Set Operations in SQL



► Set operations

- 1 UNION 3 : elements in 1 or in 3,
- 1 EXCEPT 3 : elements in 1 but not in 3 (difference),
- 1 INTERSECT 3 : elements in 1 and 3 : 2

mathematical set operations \cup , \setminus and \cap

```
SELECT name FROM Student  
UNION
```

```
SELECT name FROM TA
```

```
SELECT * FROM Student  
UNION  
SELECT * FROM TA
```

Works only if the input relations have
identical schema!

Removes duplicates (set, not multi-set)

| id | name | studies |
|----|-------|---------|
| 1 | Joe | CS |
| 2 | Jane | ITP |
| 3 | Alice | CS |
| 4 | Bob | CE |



| name |
|-------|
| Joe |
| Jane |
| Alice |
| Bob |
| Eve |
| Chris |

| TAid | name | course |
|------|-------|--------|
| 1 | Eve | DB |
| 2 | Chris | DB |
| 2 | Chris | Prog |

UNION, INTERSECT, and EXCEPT

- ▶ Corresponding multiset operations: UNION ALL, EXCEPT ALL, INTERSECT ALL
- ▶ Uses multiset semantics
 - ▶ i.e., keep duplicates

```
SELECT name FROM Student
UNION ALL
SELECT name FROM TA
```

| id | name | studies | TAid | name | course |
|----|-------|---------|------|-------|--------|
| 1 | Joe | CS | 1 | Eve | DB |
| 2 | Jane | ITP | 2 | Chris | DB |
| 3 | Alice | CS | 2 | Chris | Prog |
| 4 | Bob | CE | | | |
| | | | | Joe | |
| | | | | Jane | |
| | | | | Alice | |
| | | | | Bob | |
| | | | | Eve | |
| | | | | Chris | |
| | | | | Chris | |

- ▶ Different semantics:
 - ▶ Set: find names of all students or TAs
 - ▶ Multi-set: find names of students or TAs the number of times they occur

<https://dev.mysql.com/doc/refman/8.0/en/union.html>

<https://dev.mysql.com/doc/refman/8.0/en/intersect.html>

<https://dev.mysql.com/doc/refman/8.0/en/except.html>

Ordering of Query Results

- ▶ Use **ORDER BY** clause
 - ▶ Keyword **DESC** to see result in a descending order of values
 - ▶ 10,9,8,7,...
 - ▶ Zebra, Mouse, Elephant,...
 - ▶ Keyword **ASC** to specify ascending order explicitly (default)
 - ▶ `ORDER BY D.Dname DESC, E.Lname ASC, E.Fname ASC`
 - ▶ 1,2,3,...
 - ▶ Ape, Bear, Chimp,...



Modifications

- ▶ **SQL commands / statements** may modify the database
 - ▶ Not “query”
- ▶ Three kinds of modifications
 - ▶ insert one or more rows
 - ▶ delete one or more rows
 - ▶ update existing rows or columns
- ▶ Modifications do not return a result
 - ▶ Their effect is on the database state
 - ▶ E.g. new price value for a product, new student added, old computer removed



Inserting a Single Row

- ▶ **INSERT INTO** *table* **VALUES** (*list of values*);

```
INSERT INTO Participants  
VALUES (42835, 'ira', 'a');
```

- ▶ Optionally specify attribute names:

```
INSERT INTO  
Participants(pid, status, meetid)  
VALUES ('ira', 'a', 42835);
```

- ▶ Missing values are **NULL** or defaults

<https://dev.mysql.com/doc/refman/8.0/en/insert.html>

Example INSERTs from textbook

- ▶ Specify the relation name and a list of values for the tuple (assumes attributes and their order is known and obeyed)

```
INSERT INTO EMPLOYEE  
VALUES      ( 'Richard', 'K', 'Marini', '653298653', '1962-12-30', '98  
              Oak Forest, Katy, TX', 'M', 37000, '653298653', 4 );
```

- ▶ To insert only a subset of the values, or to specify the order (assumes remaining attributes admit NULL values or have default values):

```
INSERT INTO EMPLOYEE (Fname, Lname, Dno, Ssn)  
VALUES ('Richard', 'Marini', 4, '653298653');
```


The UPDATE Command

- ▶ Modify attribute values of one or more selected tuples
- ▶ Additional **SET** clause in the **UPDATE** command
 - ▶ Specifies attributes to be modified and new values

```
UPDATE    PROJECT
SET       Plocation = 'Bellaire', Dnum = 5
WHERE     Pnumber=10;
```

<https://dev.mysql.com/doc/refman/8.0/en/update.html>

Update example

- ▶ `UPDATE table SET attribute assignments`
`WHERE condition;`

- ▶ Move Ira to a smaller office

| userid | name | group | office |
|--------|-----------------|-------|--------|
| ira | Ira Assent | vip | Ny-357 |
| aas | Annika Schmidt | phd | NULL |
| jan | Jan Christensen | tap | Ho-017 |

```
UPDATE People
  SET office = 'Ny-343'
 WHERE userid = 'ira';
```

| userid | name | group | office |
|--------|-----------------|-------|--------|
| ira | Ira Assent | vip | Ny-343 |
| aas | Annika Schmidt | phd | NULL |
| jan | Jan Christensen | tap | Ho-017 |

The DELETE Command

- ▶ Removes tuples from a relation
 - ▶ Includes a `WHERE` clause to select the tuples to be deleted

| | |
|--------------------|-------------------------------|
| DELETE FROM | EMPLOYEE |
| WHERE | <code>Lname='Brown';</code> |
| DELETE FROM | EMPLOYEE |
| WHERE | <code>Ssn='123456789';</code> |
| DELETE FROM | EMPLOYEE |
| WHERE | <code>Dno=5;</code> |
| DELETE FROM | EMPLOYEE; |

Careful!
Practice tip: try out the condition as part of a select first to see if you get the right tuples, then issue the delete with the same condition



<https://dev.mysql.com/doc/refman/8.0/en/delete.html>

Deleting Some Rows

- ▶ **DELETE FROM *table* WHERE *condition*;**

| room | capacity |
|----------|----------|
| Ny-357 | 6 |
| Ada-333 | 26 |
| StoreAud | 286 |

- ▶ Delete Ira's office

```
DELETE FROM Rooms  
WHERE room= 'Ny-357';
```

| room | capacity |
|----------|----------|
| Ada-333 | 26 |
| StoreAud | 286 |

- ▶ Delete all offices

```
DELETE FROM Rooms;
```

| room | capacity |
|------|----------|
|------|----------|

Summary

- ▶ Intended learning outcomes
 - ▶ Be able to
 - ▶ Write simple SQL DML queries
 - ▶ Make use of advanced constructs in DML queries

Where to go from here?

- ▶ So, we can use SQL to create a database schema, to insert and manipulate data
- ▶ But, how do we ask more complicated questions?
 - ▶ How can I check if something is found in the database?
 - E.g. is there are student who has finished all courses or a professor who does not teach anything?
 - ▶ What if I want to make use of query results in another query?
 - ▶ We can express very powerful conditions in SQL using subqueries
 - SQL full power coming up!



What was this all about?

Guidelines for your own review of today's session

- ▶ A standard SQL DML query is...
 - ▶ We can specify constraints...
 - ▶ It is evaluated in the following manner...
 - ▶ A join is when...
 - ▶ We specify conditions using...
 - ▶ We order results...
- ▶ We use aggregates, do pattern matching and arithmetics...
- ▶ We can choose to remove or keep duplicates and avoid name clashes by...
- ▶ SQL handles data as multisets which are...
 - ▶ Set operations are...
- ▶ We can insert, delete and update as follows...