

# NoSQL DBMS and queries

Databases

**Ira Assent**

[ira@cs.au.dk](mailto:ira@cs.au.dk)

Data-Intensive Systems group, Department of Computer Science, Aarhus University, DK

# Intended learning outcomes

---

- ▶ Be able to
  - ▶ Describe the data model and query language in NoSQL databases
  - ▶ Discuss differences in common NoSQL systems

# Recap: Security and Authorization

- ▶ SQL injection: passing user input directly to DBMS is risky
  - ▶ Instead of expected values, may contain SQL code (parts)
  - ▶ Use variable binding and sanitizing



- ▶ Authorization to restrict data access
  - ▶ `GRANT privileges ON object TO identity`
  - ▶ `SELECT` (retrieval or read), `modify privileges` (`INSERT`, `DELETE`, `UPDATE`), `REFERENCES`, `WITH GRANT OPTION`
  - ▶ `CREATE ROLE` to group access privileges into profiles
  - ▶ Use views for fine-grained authorization for `SELECT`s
  - ▶ `REVOKE` to remove access rights again



# Recap: NoSQL data model

- ▶ NoSQL (Not Only SQL)
  - ▶ High performance, Availability, Replication, Scalability
- ▶ Willing to sacrifice
  - ▶ some (immediate) data consistency
  - ▶ powerful query languages
  - ▶ structured data storage
- ▶ Emphasize performance and flexibility over modeling power and complex querying
- ▶ Typically, **no schema** required
  - ▶ Allow **semi-structured, self-describing** data
    - ▶ JSON (JavaScript Object Notation), XML (Extensible Markup Language),...
    - ▶ Idea: description, such as tags, part of the data objects



```
<document>
  <type>Student</type>
  <name>An</name>
  <study><stname>CS</stname>
    <level>BSc</level>
  </study>
</document>
```

# Why self-describing formats in NoSQL?

---

- A. They are easier to read.
- B. They are more flexible.
- C. They are easier to normalize.
- D. They allow for better compression.

# NoSQL systems: query languages

---

- ▶ Less powerful query languages
- ▶ Many NoSQL based systems rely on simple search/read access
  - ▶ Find particular object using id
  - ▶ No need for complex query languages that express powerful conditions and combinations across tables
- ▶ Many NoSQL systems provide functions and operations via an API (application programming interface)
  - ▶ Read/write done via function/operation calls
  - ▶ CRUD (Create, Read, Update, Delete) operations
  - ▶ SCRUD (plus Search)
  - ▶ Some provide query language that corresponds to subset of SQL capabilities
    - ▶ Typically join not available
    - ▶ If needed: must be handled by application program
- ▶ Versioning
  - ▶ Some NoSQL systems provide storage of multiple versions
  - ▶ Timestamps of version creation

# Categories of NoSQL systems

---

- ▶ Four major categories of systems
  - ▶ Document-based
    - ▶ Store data as documents using e.g. JSON
    - ▶ Access using document id, other indexes
  - ▶ Key-value stores
    - ▶ Simple data model of key-value, where value can be record, document, some other (complex) data structure
  - ▶ Column-based / wide-column
    - ▶ Partition tables by column families (vertical partitioning)
    - ▶ Each column family stored in separate file
    - ▶ Typically versioning
  - ▶ Graph-based
    - ▶ Graphs as data model
    - ▶ Access data by traversing edges using path expressions

# MongoDB: document-based NoSQL

- ▶ **Document-based** also called **document store**
- ▶ Data model: collection of documents
  - ▶ Documents in a collection are typically similar, but can have different data elements (attributes)
    - ▶ Great flexibility: if you feel like adding an attribute, just put in a new element in the document in question
  - ▶ Documents are self-describing, typically BSON based (binary JSON) format



```
"document": {  
  "type": "Student Record",  
  "name": "Anne Christensen",  
  "study": "Computer Science",  
  {  
    "level": "Bachelor"  
  }  
}
```





# MongoDB: document-based NoSQL

---

## COMPANY database example

- ▶ **Create** a collection called project to hold PROJECT objects

```
db.createCollection("project", {capped:  
true, size: 1310720, max: 500})
```

- ▶ **Capped:** limit on storage (size) and number of documents (max)

```
db.createCollection("worker", {capped:  
true, size: 5242880 max: 2000})
```

- ▶ Each document in a collection has **automatically indexed unique identifier** called `_id`



# MongoDB CRUD operations

---

- ▶ CRUD: create, read, update, delete

- ▶ **Create and insert** documents into collection using

```
db.<collection_name>.insert(<document_name(s)>)
```

- ▶ Textbook example

```
db.project.insert( { _id: "P1", Pname: "ProductX", Plocation: "Bellaire" } )
```

```
db.worker.insert( [ { _id: "W1", Ename: "John Smith", ProjectId: "P1", Hours: 32.5 },  
                    { _id: "W2", Ename: "Joyce English", ProjectId: "P1",  
                      Hours: 20.0 } ] )
```

- ▶ **Delete:** `db.<collection_name>.remove(<condition>)`
    - ▶ **Read:** `db.<collection_name>.find(<condition>)`
    - ▶ General Boolean conditions evaluating to true are selected



# Document-based NoSQL and MongoDB

---

- ▶ Collection **does not have schema**
  - ▶ Structure of data fields in documents
    - ▶ **Normalized**
      - Decompose into documents of similar structure and content
        - References between documents
          - ▶ Store ids of other documents in the document
    - ▶ **Denormalized**
      - All information in one document



# Denormalized pattern

- ▶ Textbook example:  
store workers John,  
Joyce in ProductX  
document
- ▶ Information about  
workers and projects  
stored together
  - ▶ In relation model:  
corresponds to storing  
also the workers in the  
project table
  - ▶ **denormalized**
  - ▶ [...] notation: array
- ▶ information about  
workers **embedded** in  
project document, no  
need for separate  
worker collection

project document with an array of embedded workers:

```
{
  _id:          "P1",
  Pname:        "ProductX",
  Plocation:    "Bellaire",
  Workers: [
    { Ename: "John Smith",
      Hours: 32.5
    },
    { Ename: "Joyce English",
      Hours: 20.0
    }
  ]
};
```

# Embedded array of references

- ▶ **Alternative:** project document with an embedded array of worker ids:  
**embed**  
**only**  
**references,**  
here to  
workers via  
their id
- ```
{
  _id:          "P1",
  Pname:       "ProductX",
  Plocation:   "Bellaire",
  WorkerIds:   [ "W1", "W2" ]
}

{ _id:          "W1",
  Ename:       "John Smith",
  Hours:      32.5
}

{ _id:          "W2",
  Ename:       "Joyce English",
  Hours:      20.0
}
```

# What's this?

- A. Fully normalized representation of projects and workers
  - B. Normalized representation of projects and workers, but not for M:N relationships
  - C. Denormalized representation of projects and workers
  - D. Denormalized representation of workers embedded in projects
- ```

{
  _id: "P1",
  Pname: "ProductX",
  Plocation: "Bellaire"
}
{
  _id: "W1",
  Ename: "John Smith",
  ProjectId: "P1",
  Hours: 32.5
}
{
  _id: "W2",
  Ename: "Joyce English",
  ProjectId: "P1",
  Hours: 20.0
}
    
```

# Normalized documents

- ▶ Fully **normalized** for many-to-many requires three collections project, employee, works\_on (as seen in relational model before)
- ▶ Here: employee working on several projects represented by several worker documents with different ids
- ▶ **redundancy** normalized project and worker documents (not a fully normalized design for M:N relationships):

```
{
  _id:          "P1",
  Pname:        "ProductX",
  Plocation:    "Bellaire"
}
{
  _id:          "W1",
  Ename:        "John Smith",
  ProjectId:    "P1",
  Hours:        32.5
}
{
  _id:          "W2",
  Ename:        "Joyce English",
  ProjectId:    "P1",
  Hours:        20.0
}
```

# Voldemort

---

- ▶ Voldemort open source project builds on Amazon DynamoDB

- ▶ Data model:

- ▶ Collections of self-describing items
  - ▶ NOT like relational model, without schema
- ▶ Items are attribute-value pairs
  - ▶ Item can also be JSON file



- ▶ **Read** the entry with key 1234

- ▶ `get(1234) => {"name": "ann",  
"email": "ann.jensen@cs.au.dk"}`

- ▶ **Write** the entry with key 1234

- ▶ `put(1234, {"name": "ann jensen",  
"email": "ann.jensen@cs.au.dk"})`

- ▶ **Delete** the entry with key 1234

- ▶ `Delete(1234)` <https://www.project-voldemort.com/voldemort/design.html>



# Key-Value Storage

---




- ▶ Only very simple data model: key-value data access
  - ▶ keys and values can be complex compound objects (including JSON)
- ▶ High performance, availability and scalability via distributed storage system
- ▶ No expressive query language, but set of operations to be used in application programs
- ▶ Pros
  - ▶ efficient queries, very predictable performance
  - ▶ easy to distribute across a number of sites (distributed database)
  - ▶ no object-relational impedance mismatch
- ▶ Cons
  - ▶ no complex queries
  - ▶ all joins must be done in program code
  - ▶ no foreign key constraints
  - ▶ no triggers

# Scaling Voldemort storage




- ▶ Key-value data access across distributed database supported by consistent hashing
  - ▶ Map the key value to a hash value
  - ▶ Range of hash values stored by some site (or several for replication)
- ▶ Hash function  $h(\text{key})$  applied to key of (key, value) pair
  - ▶ E.g.  $h(\text{key}) = \text{value} \bmod 5$ 
    - ▶  $h(3)=3, h(7)=2, h(10)=0$
  - ▶ E.g. A: range 1, B: range 2, C: range 3
    - ▶ Ranges could be hash values 1-3, 4, 5-0
    - ▶ Ideally, each server gets the same workload and amount of data





# How do we scale to more data?

- A. Add a range and a site
- B. Add a site
- C. Redistribute ranges to sites
- D. Add rings



# Horizontal scalability

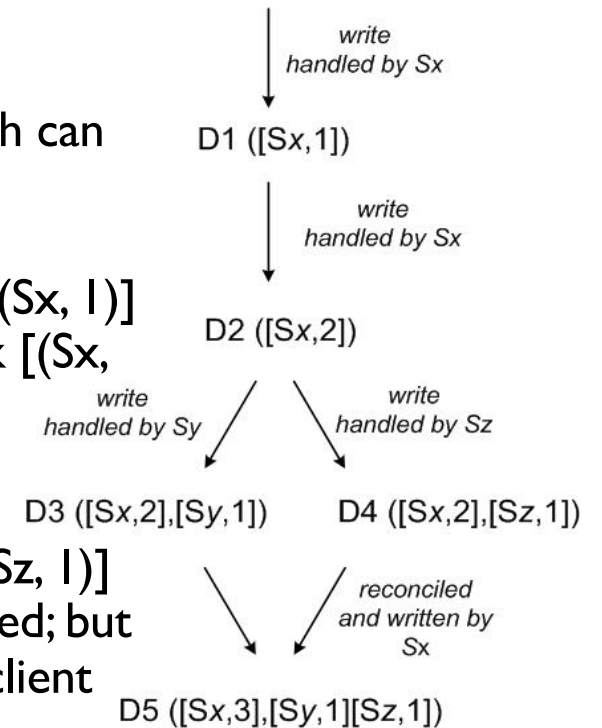
- ▶ **Easy horizontal scalability**
  - ▶ Add node to ring
    - ▶ E.g. add D
      - ▶ Items in range 4 are moved to D from B and C (ranges 2 and 3 reduced)





# Consistency and versioning

- ▶ Voldemort allows concurrent writes
  - ▶ Different values might be associated with the same key at different nodes when items replicated
  - ▶ When reading, achieve consistency by "versioning and read repair"
    - ▶ Each write associated with vector clock value
      - If system can reconcile different read values, do so
      - Otherwise, pass on several values to application which can reconcile based on application semantics
- ▶ Example:
  - ▶ Node  $S_x$  writes object  $D1$  and produces associated clock  $[(S_x, 1)]$
  - ▶  $S_x$  then handles it again and produces version  $D2$  and clock  $[(S_x, 2)]$
  - ▶  $D2$  overwrites  $D1$ , but  $D1$  copies may exist still
  - ▶ Node  $S_y$  updates  $D2$  to  $D3$  with clocks  $[(S_x, 2), (S_y, 1)]$
  - ▶  $S_z$  updates  $D2$  to  $D4$  (unaware of  $D3$ ) with clock  $[(S_x, 2), (S_z, 1)]$
  - ▶  $S_x$  finds  $D1$  and  $D2$  are overwritten, can be garbage collected; but for  $D3$  and  $D4$  finds no relation, so both must be kept for client (upon a read) for semantic reconciliation



# Hbase



- ▶ Column-based / wide column system, similar to Google's BigTable
- ▶ Apache Hbase typically uses HDFS (Hadoop Distributed File System) for storage
- ▶ Sparse multidimensional distributed persistent sorted map
  - ▶ Map is a collection of (key, value) pairs, key is mapped to value
  - ▶ Main difference is the key: here, multidimensional, typically a combination of table name, row key, column, timestamp
  - ▶ Column typically consists of column family and column qualifier

<https://hbase.apache.org/>

# Hbase examples

- ▶ Table name followed by the names of the column families
  - ▶ E.g. Name
- ▶ Column qualifiers may group related columns for storage purposes when data is created (self-describing; not defined at creation time)
  - ▶ E.g. Fname

## Some Hbase basic CRUD operations:

Creating a table: create <tablename>, <column family>, <column family>, ...

Inserting Data: put <tablename>, <rowid>, <column family>:<column qualifier>, <value>

Reading Data (all data in a table): scan <tablename>

Retrieve Data (one item): get <tablename>,<rowid>

## creating a table:

```
create 'EMPLOYEE', 'Name', 'Address', 'Details'
```

## inserting some row data in the EMPLOYEE table:



```
put 'EMPLOYEE', 'row1', 'Name:Fname', 'John'
put 'EMPLOYEE', 'row1', 'Name:Lname', 'Smith'
put 'EMPLOYEE', 'row1', 'Name:Nickname', 'Johnny'
put 'EMPLOYEE', 'row1', 'Details:Job', 'Engineer'
put 'EMPLOYEE', 'row1', 'Details:Review', 'Good'
put 'EMPLOYEE', 'row2', 'Name:Fname', 'Alicia'
put 'EMPLOYEE', 'row2', 'Name:Lname', 'Zelaya'
put 'EMPLOYEE', 'row2', 'Name:MName', 'Jennifer'
put 'EMPLOYEE', 'row2', 'Details:Job', 'DBA'
put 'EMPLOYEE', 'row2', 'Details:Supervisor', 'James Borg'
put 'EMPLOYEE', 'row3', 'Name:Fname', 'James'
put 'EMPLOYEE', 'row3', 'Name:Minit', 'E'
put 'EMPLOYEE', 'row3', 'Name:Lname', 'Borg'
put 'EMPLOYEE', 'row3', 'Name:Suffix', 'Jr.'
put 'EMPLOYEE', 'row3', 'Details:Job', 'CEO'
put 'EMPLOYEE', 'row3', 'Details:Salary', '1,000,000'
```





# Neo4j

- ▶ Example graph oriented NoSQL system
- ▶ Data model: graph
  - ▶ Collection of vertices (nodes) and edges
    - ▶ Can be labeled to indicate the types of entities and relationships they represent
- ▶ Open source, Java
- ▶ High-level query language Cypher
  - ▶ Declarative commands for creating, modifying, and finding nodes and relationships



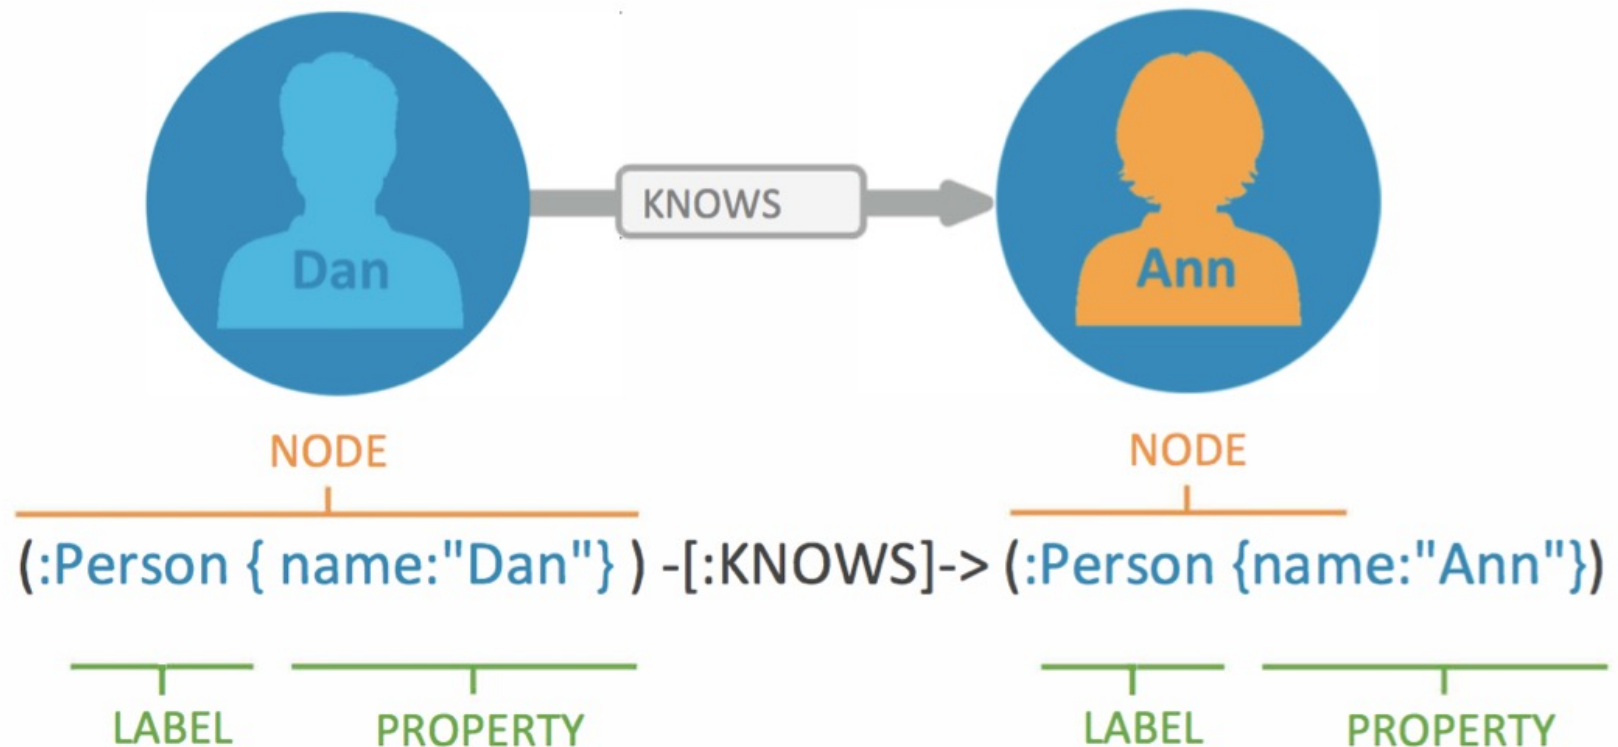
<https://neo4j.com/docs/cypher-manual/current/>

<https://neo4j.com/docs/getting-started/current/cypher-intro/patterns/#cypher-intro-patterns>



# Neo4j Graph Query Example

- ▶ Start from edge Knows between nodes with label Person and property Dan, Ann, resp.



# Neo4j Graph Query Example: Cypher

- ▶ Generalize by introducing Alias to second property of label Person
  - ▶ This generalized pattern is used in `MATCH` to define a query



# Graph Queries: CRUD


---

- ▶ **Create Read Update Delete** operations in Cypher query language
  - ▶ Similar to insert, select, update, delete in relational model
  - ▶ Examples:
    - ▶ Create new node with properties P
    - ▶ Add edge from v1 to v2 (plus some properties P)
    - ▶ Add property { Name :Value } to node v or to edge e
    - ▶ Add a new node, and then edges from it to other
    - ▶ Find nodes/edges with property P { Name :Value }
    - ▶ Find edges with a specific label

Operations on nodes and edges, and the patterns they are part of

# Example Cypher query

- ▶ Find nodes/edges with property P { Name :Value }
- ▶ Find edges with a specific label
- ▶ In Cypher:
- ▶ **MATCH** pattern **WHERE** predicate **RETURN** expression



```
MATCH (node1:Person) -[:KNOWS]->
      (node2:Person)
      (node1) -[:LIVES_IN]->(node3:City)
      (node2) -[:LIVES_IN]->(node3)
WHERE node1.propertyA = {value1}
RETURN node2.propertyA, node2.propertyB
```

# Creating nodes for COMPANY

- ▶ **EMPLOYEE, DEPARTMENT, etc. are node labels**
    - ▶ **Multiple labels possible**
  - ▶ **{...} contain properties**
- ```

CREATE (e1: EMPLOYEE, {Empid: '1', Lname: 'Smith', Fname: 'John', Minit: 'B'})
CREATE (e2: EMPLOYEE, {Empid: '2', Lname: 'Wong', Fname: 'Franklin'})
CREATE (e3: EMPLOYEE, {Empid: '3', Lname: 'Zelaya', Fname: 'Alicia'})
CREATE (e4: EMPLOYEE, {Empid: '4', Lname: 'Wallace', Fname: 'Jennifer', Minit: 'S'})
...
CREATE (d1: DEPARTMENT, {Dno: '5', Dname: 'Research'})
CREATE (d2: DEPARTMENT, {Dno: '4', Dname: 'Administration'})
...
CREATE (p1: PROJECT, {Pno: '1', Pname: 'ProductX'})
CREATE (p2: PROJECT, {Pno: '2', Pname: 'ProductY'})
CREATE (p3: PROJECT, {Pno: '10', Pname: 'Computerization'})
CREATE (p4: PROJECT, {Pno: '20', Pname: 'Reorganization'})
...
CREATE (loc1: LOCATION, {Lname: 'Houston'})
CREATE (loc2: LOCATION, {Lname: 'Stafford'})
CREATE (loc3: LOCATION, {Lname: 'Bellaire'})
CREATE (loc4: LOCATION, {Lname: 'Sugarland'})
...

```



# Creating relationships

## ▶ CREATE syntax

- ▶ → indicate direction of relationships
- ▶ E.g. e1 works for d1
- ▶ Relationship types are labels WorksFor, Manager,...

(d1: DEPARTMENT, {Dno: '5', Dname: 'Research'})

WorksFor

e1: EMPLOYEE, {Empid: '1', Lname: 'Smith', Fname: 'John', Minit: 'B'})

CREATE (e1) - [ : WorksFor ] -> (d1)

CREATE (e3) - [ : WorksFor ] -> (d2)

...

CREATE (d1) - [ : Manager ] -> (e2)

CREATE (d2) - [ : Manager ] -> (e4)

...

CREATE (d1) - [ : LocatedIn ] -> (loc1)

CREATE (d1) - [ : LocatedIn ] -> (loc3)

CREATE (d1) - [ : LocatedIn ] -> (loc4)

CREATE (d2) - [ : LocatedIn ] -> (loc2)

...

CREATE (e1) - [ : WorksOn, {Hours: '32.5'} ] -> (p1)

CREATE (e1) - [ : WorksOn, {Hours: '7.5'} ] -> (p2)

CREATE (e2) - [ : WorksOn, {Hours: '10.0'} ] -> (p1)

CREATE (e2) - [ : WorksOn, {Hours: 10.0} ] -> (p2)

CREATE (e2) - [ : WorksOn, {Hours: '10.0'} ] -> (p3)

CREATE (e2) - [ : WorksOn, {Hours: 10.0} ] -> (p4)



RELATIONSHIPS



# Queries

---

- ▶ Cypher queries made up of clauses
  - ▶ Result of one clause may be used as input to another
  - ▶ Similar to what we had seen for SQL (tables in, tables out)

## **Basic simplified syntax of some common Cypher clauses:**

Finding nodes and relationships that match a pattern: `MATCH <pattern>`

Specifying aggregates and other query variables: `WITH <specifications>`

Specifying conditions on the data to be retrieved: `WHERE <condition>`

Specifying the data to be returned: `RETURN <data>`

Ordering the data to be returned: `ORDER BY <data>`

Limiting the number of returned data items: `LIMIT <max number>`

Creating nodes: `CREATE <node, optional labels and properties>`

Creating relationships: `CREATE <relationship, relationship type and optional properties>`

Deletion: `DELETE <nodes or relationships>`

Specifying property values and labels: `SET <property values and labels>`

Removing property values and labels: `REMOVE <property values and labels>`

# What could a relational equivalent be?

MATCH(e: EMPLOYEE {Empid:'2'}) – [w :WorksOn] → (p)

RETURN (e.Ename, w.Hours, p.Pname)

- A. SELECT Ename,Hours,Pname FROM EMPLOYEE WHERE Empid='2';
- B. SELECT Ename,Hours,Pname FROM EMPLOYEE,WORKSON WHERE Empid='2';
- C. SELECT Ename,Hours,Pname FROM EMPLOYEE NATURAL JOIN WORKSON WHERE Empid='2';
- D. SELECT Ename,Hours,Pname FROM EMPLOYEE NATURAL JOIN WORKSON NATURAL JOIN Projects WHERE Empid='2';



# Query examples

- ▶ **MATCH** specifies a pattern and query variables (d, loc)
- ▶ **RETURN** specifies query result

## Examples of simple Cypher queries:

1. `MATCH (d : DEPARTMENT {Dno: '5'}) - [ : LocatedIn ] → (loc)  
RETURN d.Dname , loc.Lname`
2. `MATCH (e: EMPLOYEE {Empid: '2'}) - [ w: WorksOn ] → (p)  
RETURN e.Ename , w.Hours, p.Pname`
3. `MATCH (e) - [ w: WorksOn ] → (p: PROJECT {Pno: 2})  
RETURN p.Pname, e.Ename , w.Hours`
4. `MATCH (e) - [ w: WorksOn ] → (p)  
RETURN e.Ename , w.Hours, p.Pname  
ORDER BY e.Ename`
5. `MATCH (e) - [ w: WorksOn ] → (p)  
RETURN e.Ename , w.Hours, p.Pname  
ORDER BY e.Ename  
LIMIT 10`
6. `MATCH (e) - [ w: WorksOn ] → (p)  
WITH e, COUNT(p) AS numOfprojs  
WHERE numOfprojs > 2  
RETURN e.Ename , numOfprojs  
ORDER BY numOfprojs`
7. `MATCH (e) - [ w: WorksOn ] → (p)  
RETURN e , w, p  
ORDER BY e.Ename  
LIMIT 10`
8. `MATCH (e: EMPLOYEE {Empid: '2'})  
SET e.Job = 'Engineer'`

# Graph Query Example: OrientDB SQL



- ▶ **OrientDB:** Open source database written in Java (now owned by SAP)
  - ▶ multimodal, i.e., supports data models graph, document, key value and object

**/\* A and B are friends \*/**

```
MATCH {class:Person, as:A} -FriendOf- {class:Person,  
as:B},
```

**/\* C is a friend of both A and B \*/**

```
{as:A} -FriendOf- {as:C} -FriendOf- {as:B},
```

**/\* A has a cat \*/**

```
{as:C} -HasA-> {class:Cat},
```

**/\* B works at AU \*/**

```
{as:B} -WorksAt-> {class:Company,  
where: (name="AU") },
```

**/\* return the names of A and B with aliases name, friendName\*/**

```
RETURN A.name as name, B.name as friendName
```

<https://orientdb.com/nosql/pattern-matching-with-orientdb/>

# Summary

---

- ▶ Intended learning outcomes
- ▶ Be able to
  - ▶ Describe the data model and query language in NoSQL databases
  - ▶ Discuss differences in common NoSQL systems

# Where to go from here?

---

- ▶ We have seen NoSQL as an approach for high performance and availability of very large databases where immediate consistency and complex query languages are not the focus
- ▶ We will now turn to handling text documents, in particular from the web
  - ▶ Information Retrieval and Web



# What was this all about?

Guidelines for your own review of today's session

---

- ▶ NoSQL systems treat consistency as...
  - ▶ Which is different from relational DBMS where...
- ▶ NoSQL data models...
- ▶ Self-describing formats...
  - ▶ E.g. we have...
- ▶ NoSQL query languages are....
  - ▶ They typically provide operations: ...
- ▶ Some types of NoSQL systems are...
- ▶ MongoDB uses the following....
- ▶ Neo4j instead models data as...
  - ▶ A query is expressed by...