

Today:

- **break planned ~13-13.10**
- **end planned ~ 13.50**

(I have to give a talk after the lecture)

Concurrency Control

Databases, Aarhus University

Ira Assent

Intended learning outcomes

- ▶ Be able to
 - ▶ Determine serializable schedules
 - ▶ Apply the two-phase locking protocol

Welcome to IADB students

- ▶ Instructor: Ira Assent
- ▶ “Super” TAs Pernille Matthews and Cheng Huang
- ▶ TAs for exercises classes and at **study café**
- ▶ **Lectures:** Tuesdays 12-14, Thursdays 10-12
 - ▶ **TA sessions** Tue afternoon → Tue morning
- ▶ Weekly **handins** and multiple choice **exam** in June
 - ▶ Handins count for 20% of grade, exam for 80%
 - ▶ Need to sign up for handin groups (contact TA regarding enrollment)
- ▶ **Forum:** make good use of it: your questions here benefit everyone
- ▶ Please check “the to success” on homepage
- ▶ DBS pre-requisite
 - ▶ Access to earlier material from this year if you need to brush up
- ▶ See also details on brightspace page
- ▶ And for the rest of you: Welcome back after Easter!



Recap: Transactions

- ▶ Transaction
 - ▶ **Logical unit** of database processing
 - ▶ includes one or more operations on the database
- ▶ Basic operations are read and write
 - ▶ read_item(X)
 - ▶ Reads a database item named X
 - ▶ E.g. into a program variable
 - ▶ write_item(X)
 - ▶ Writes the value of database item X
 - ▶ We do not need any other information to understand the impact of transactions on the database state
 - ▶ I.e., which value a database item has, and which transaction has seen which value of a database item
- ▶ BEGIN **or** START TRANSACTION, COMMIT, ROLLBACK

<https://dev.mysql.com/doc/refman/8.0/en/commit.html>

<https://dev.mysql.com/doc/refman/8.0/en/sql-transactional-statements.html>

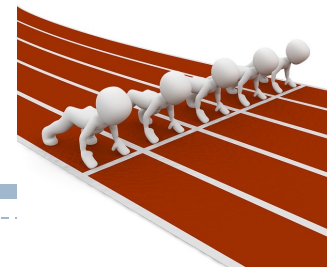


Recap:ACID



- ▶ DBMS supports **ACID** transactions
 - ▶ **Atomic** : Either the whole transaction process is done or none is
 - ▶ Any transaction does all its work on the database or none at all (atomicity)
 - so that there are no partial updates to the data or the like
 - ▶ **Consistent** : Database constraints are preserved
 - ▶ Any transaction that starts on a consistent database leaves it again in a consistent state; temporary inconsistency while it is still working is permitted;
 - so that there are no data integrity issues or the like
 - ▶ **Isolated** : It appears to user as if only one transaction process executes at a time
 - ▶ No transaction interferes with another
 - So that transactions do not produce conflicting results; looks like first one of the transactions has the database to itself until done, then the other
 - ▶ **Durable** : Effects of a process do not get lost if the system crashes
 - ▶ So that the changes made by a transaction actually persist and do not disappear
 - E.g. data is not actually written to the disk or the like

Concurrency control and schedules



- ▶ Goal of concurrency control: support more than one user / query at the same time
 - ▶ Sequential / serial: simply queue → slow
- ▶ Concurrent: allow access to database at the same time in interleaved fashion, but ensure same result
- ▶ **Schedule**: sequence of operations from one or more transactions
 - ▶ Operations considered: read(A), write (A) (start, commit, end, abort)
 - ▶ In concurrent transactions, the operations are interleaved

ACID: What is concurrency control mostly concerned with?

- A. A, I
- B. A, C
- C. A, D
- D. I, C
- E. I, D
- F. C, D



Types of schedules

▶ **Serial** schedule:

- ▶ Schedule S serial if, for every transaction T in S, all operations of T executed consecutively in S
 - ▶ Otherwise, **nonserial** schedule
- ▶ Simplest case: each transaction gets its turn to have the database by itself, while others wait

▶ **Serializable** schedule:

- ▶ Schedule S serializable if **equivalent** to some serial schedule of same transactions
- ▶ Typically, operations of different transactions interweaved, but without creating issues
- ▶ These are the kinds of schedules we want to create
 - ▶ Fast, multi-user access and **correct** result

Example

- ▶ Database with two items, X and Y , program variables x, y

- ▶ Only criterion for correctness: $X=Y$

- ▶ Two concurrent transactions for programs

T_1 : $X := X + 1$

$Y := Y + 1$

T_2 : $X := 2 * X$

$Y := 2 * Y$

- ▶ Initially, $X=10$ and $Y=10$
- ▶ Serial schedule to the right:
 - ▶ execute first T_2 entirely, then T_1 entirely
 - ▶ Resulting database state:
 - ▶ $X = 21, Y = 21, X = Y$
 - ▶ Alternatively, first T_1 , then T_2 : $X=Y=22$

T_1	T_2
	read_item(X); X:=2*X; write_item(X); read_item(Y); Y:=2*Y; write_item(Y);
read_item(X); X:=X+1; write_item(X); read_item(Y); Y:=Y+1; write_item(Y);	

Motivation for serializability

T_1	T_2
read_item(X); $X := X + 1$; write_item(X);	
	read_item(X); $X := 2 * X$; write_item(X);
	read_item(Y); $Y := 2 * Y$; write_item(Y);
read_item(Y); $Y := Y + 1$; write_item(Y);	

- ▶ Alternative concurrent schedule “incorrect”: $X=22, Y=21, X \neq Y$
- ▶ Concurrent execution of transactions **correct** if equivalent to some serial execution of those transactions
- ▶ For transactions T_1 and T_2 there are only two serial schedules: T_1, T_2 or T_2, T_1 and the result of this schedule is different from both of these
 - ▶ $X=Y=21$ or $X=Y=22$
 - ▶ Why? Because both T_1 and T_2 write X

Defining equivalence of schedules

- ▶ When are two schedules **equivalent**?
 - ▶ In particular, we would like to understand if an interweaved schedule is equivalent to a serial schedule
- ▶ Two schedules are **result equivalent** if they produce the same final state of the database
 - ▶ This can happen by “accident” – state of the database is the same, but the operations differ → cannot be used to find the right schedules
- ▶ Instead, look at the operations in the transactions
 - ▶ Do they interfere with one another?
 - ▶ Notion of **conflict**



Determining Serializability

- ▶ TAs T_i and T_j **conflict** iff there exists some item X , accessed by both T_i and T_j , and at least one of them wrote X
 - ▶ Intuitively, conflict between TAs forces an execution order between them
 - ▶ If there are no writes, the database state does not change, so no issue if both of them read it
- ▶ Let I and J be consecutive instructions by two different transactions within schedule S
 - ▶ If I and J do not conflict, may swap order to produce new schedule S'
 - ▶ S and S' **conflict equivalent**
- ▶ Schedule **conflict serializable** if conflict equivalent to a serial schedule



Serializability

This schedule with an initial database satisfying $X = Y$ is

- A. Equivalent to a serial execution of T_3, T_4
- B. Equivalent to a serial execution of T_4, T_3
- C. Equivalent to a serial execution of T_3, T_4 and T_4, T_3
- D. Not equivalent to a serial execution of T_3, T_4 or T_4, T_3

T_3	T_4
read_item(X);	
$X := X + 1$;	
	read_item(X);
write_item(X);	
	$X := 2 * X$;
	write_item(X);
	read_item(Y);
	$Y := 2 * Y$;
read_item(Y);	
$Y := Y + 1$;	
write_item(Y);	
	write_item(Y);



Possible Transaction Conflicts

▶ Write/Read conflict:

- ▶ T_2 must be executed after T_1 , as T_2 reads value provided by T_1

T_1	T_2
write_item(X);	read_item(X);

▶ Read/Write conflict:

- ▶ T_2 must be executed after T_1 , as T_2 writes a new value after T_1 reads the old value

T_1	T_2
read_item(X);	write_item(X);

▶ Write/Write conflict:

- ▶ T_2 must be executed after T_1 , as T_2 overwrites value created by T_1

T_1	T_2
write_item(X);	write_item(X);

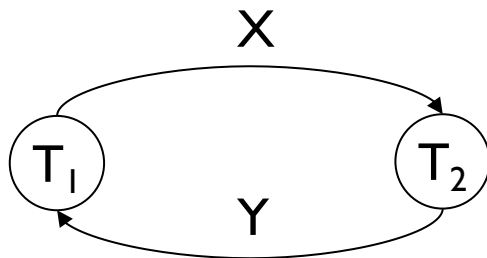
▶ No conflict:

- ▶ no implied execution order of T_1 and T_2 as both read same value of X

T_1	T_2
read_item(X);	read_item(X);

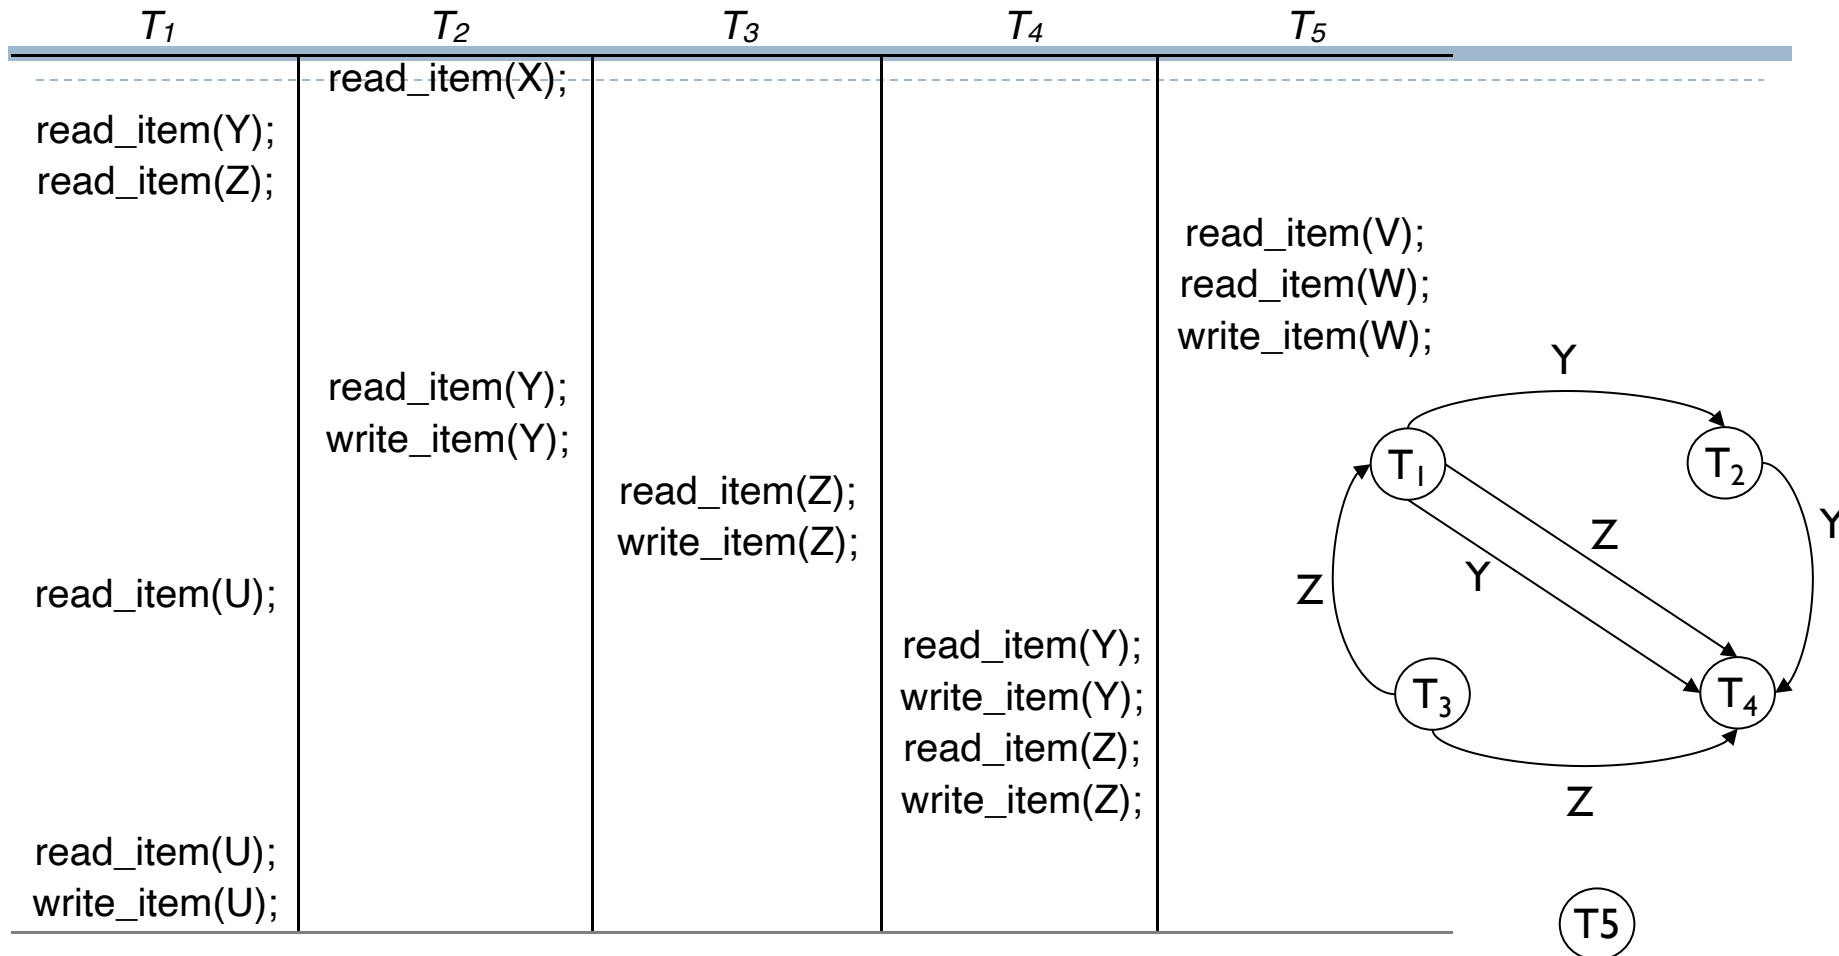
Precedence Graph (serialization graph)

- ▶ Directed graph depicting conflicts in schedule
- ▶ Each transaction is a node
- ▶ edge labeled X from T_i to T_j if T_i conflicts on X (and is before) T_j



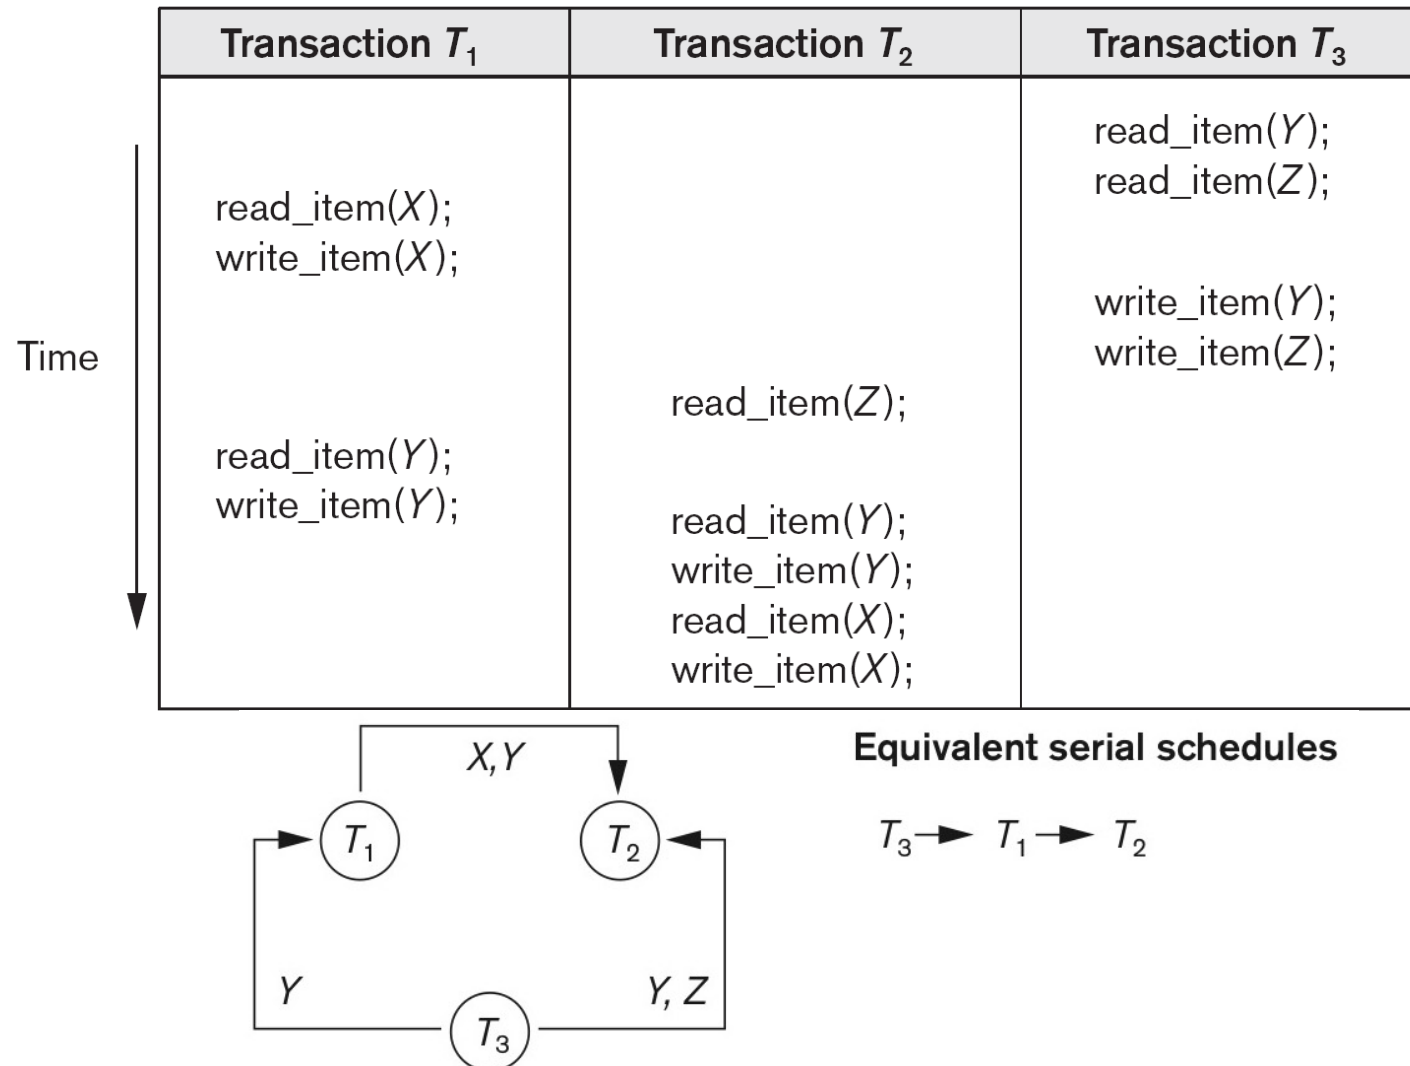
T_1	T_2
read_item(X); $X := X + 1$; write_item(X);	read_item(X); $X := 2 * X$; write_item(X); read_item(Y); $Y := 2 * Y$; write_item(Y);
read_item(Y); $Y := Y + 1$; write_item(Y);	

Is this precedence graph correct?



- A. No, T_5 is not connected
- B. No, the label from T_2 to T_4 should be Z
- C. No, T_2 should precede T_1
- D. No, T_1 should precede T_3
- E. No, arrows point the wrong way
- F. Yes

Textbook Example of Serializability Testing



Another Example of Serializability Testing

	Transaction T_1	Transaction T_2	Transaction T_3
Time ↓	read_item(X); write_item(X);	read_item(Z); read_item(Y); write_item(Y);	read_item(Y); read_item(Z);
	read_item(Y); write_item(Y);	read_item(X); write_item(X);	write_item(Y); write_item(Z);

Equivalent serial schedules

None

Reason

Cycle $X(T_1 \rightarrow T_2), Y(T_2 \rightarrow T_1)$

Cycle $X(T_1 \rightarrow T_2), YZ(T_2 \rightarrow T_3), Y(T_3 \rightarrow T_1)$

Testing Serializability


- ▶ Test: schedule is (conflict) serializable if its precedence graph is acyclic
- ▶ Topological sorting gives a serialization order
- ▶ Need to guarantee serializability
 - ▶ An inefficient strategy
 - ▶ Generate a schedule, build the precedence graph, and test for cycle
 - ▶ If cycle found, need to generate another schedule, ...



Determining serializability

- ▶ Serializability hard to check
 - ▶ Interleaving of operations in operating system through some scheduler
 - ▶ Difficult to determine beforehand how operations in schedule will be interleaved

Practical approach:

- ▶ Devise **protocols** (methods) to ensure serializability
- ▶ Not possible to determine when schedule begins and ends
 - ▶ “the database never sleeps” – endless stream of transactions 
 - ▶ Reduce the problem of checking the whole schedule to checking only a **committed projection** of the schedule
 - ▶ i.e. operations from only the committed transactions
- ▶ Current approach used in most DBMSs:
 - ▶ Use of locks with two phase locking

Locks for concurrency control



- ▶ Locking is an operation which secures
 - ▶ (a) permission to read and/or
 - ▶ (b) permission to write a data item for a transaction
 - ▶ **Lock (X)**: data item X is locked in behalf of the requesting transaction, so it obtains access permission
- ▶ Unlocking is an operation which removes these permissions from the data item
 - ▶ **Unlock (X)**: data item X is made available to all other transactions
- ▶ Lock and Unlock are atomic operations
- ▶ Transaction must be **well-formed**:
 - ▶ It must lock the data item before it reads or writes to it
 - ▶ It must not lock an already locked data item and it must not try to unlock a free data item
- ▶ Lock Manager maintains lock table with information on (at least) data items with locks currently granted, queue for waiting transactions

Lock modes

- ▶ Two **locks modes**:
- ▶ Shared mode: **shared lock (X)** or **read lock (X)**
 - ▶ More than one transaction can apply shared lock on X for reading its value, but no write lock can be applied on X by any other transaction
- ▶ Exclusive mode: **exclusive lock (X)** or **write lock (X)**
 - ▶ Only one write lock on X can exist at any time and no shared lock can be applied by any other transaction on X
- ▶ Conflict matrix
 - ▶ Same as for conflict serializability

	Read	Write
Read	Y	N
Write	N	N



The Two-Phase Protocol

- ▶ Phase 1 (*growing phase*):
 - ▶ Transaction may request locks
 - ▶ Transaction may not release locks
- ▶ Phase 2 (*shrinking phase*):
 - ▶ Transaction may not request locks
 - ▶ Transaction may release locks
- ▶ When the first lock is released, the transaction moves from phase 1 to phase 2



Two-Phase With Lock Conversion

- ▶ First phase
 - ▶ Can acquire a shared lock on item X
 - ▶ Can acquire an exclusive lock on item X
 - ▶ Can **convert (upgrade)** a shared lock on X to an exclusive lock on X
- ▶ Second phase
 - ▶ Can release a shared lock
 - ▶ Can release an exclusive lock
 - ▶ Can **convert (downgrade)** an exclusive lock to a shared lock
- ▶ This protocol assures serializability



Practical Two-Phase Protocol (2PL)



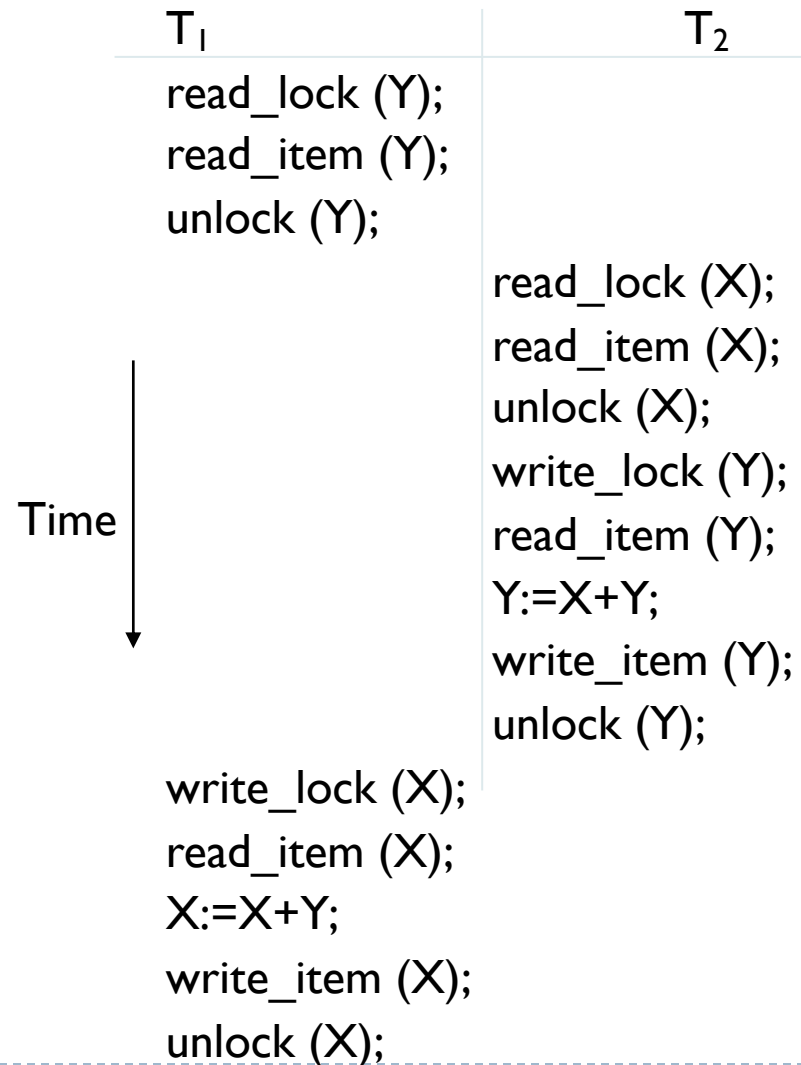
- ▶ Transaction issues standard read/write instructions
- ▶ System manages protocol, including lock operations

read X: if T has a lock on X
 perform read
 else wait until no other transaction has write_lock(X)
 grant read_lock(X) to T
 perform read

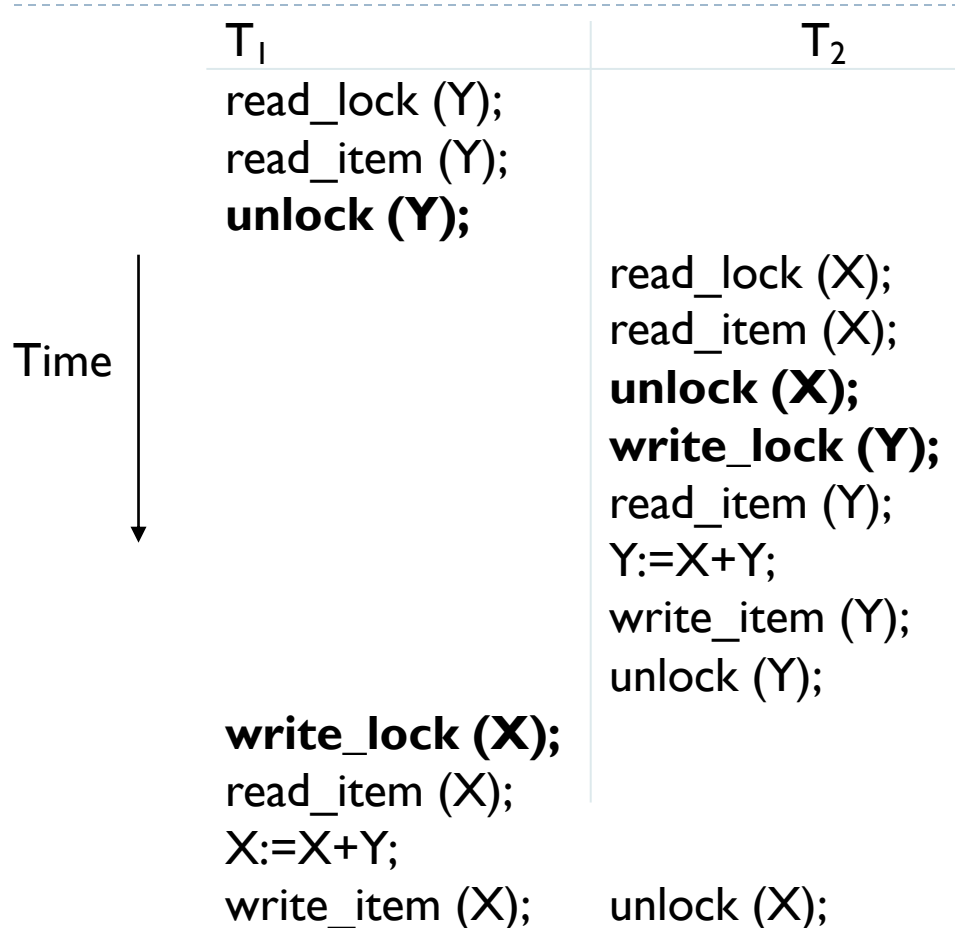
write X: if T has a write_lock(X)
 perform write
 else wait until no other transaction has a lock on X
 if T has read_lock(X)
 convert it to write_lock(X)
 else grant it write_lock(X)
 perform write

2PL?

- A. Yes, follows 2PL
- B. No, phases not separate
- C. Yes, if lock conversion is applied
- D. No, there is a conflict



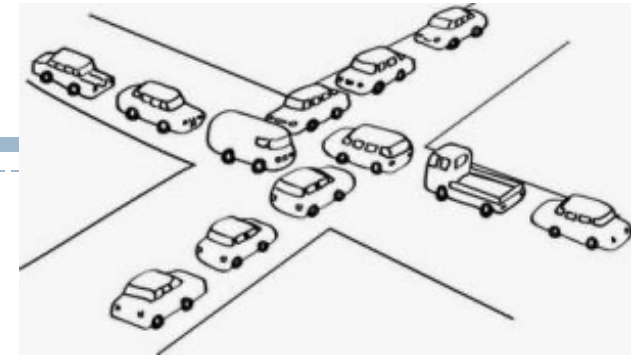
Textbook locking example



This is **NOT** 2PL, as e.g. T_1 unlocks Y and then locks X
→ would not be allowed to run



Textbook example: 2PL



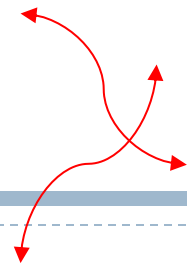
T_1	T_2
read_lock (Y);	read_lock (X);
read_item (Y);	read_item (X);
	write_lock (Y);
write_lock (X);	
unlock (Y);	
read_item (X);	unlock (X);
$X := X + Y;$	read_item (Y);
write_item (X);	$Y := X + Y;$
unlock (X);	write_item (Y);
	unlock (Y);

Red arrows indicate dependencies: one from **read_lock (Y)** to **write_lock (Y)**, and another from **write_lock (X)** to **unlock (X)**.

► Now 2PL policy followed, but **deadlock**

- T_1 cannot get **write_lock(X)** as T_2 has **read_lock(X)** and T_2 cannot get **write_lock(Y)** as T_1 has **read_lock(Y)**

Deadlocks



▶ **Deadlock**

- ▶ A cycle of transactions waiting for one another's unlock (cycle wait)

▶ **Deadlock prevention**

- ▶ A transaction locks all data items it refers to before it begins execution
 - ▶ prevents deadlock since a transaction never waits for a data item
 - ▶ Impractical: usually necessary data items not fully known

▶ **Deadlock detection and resolution**

- ▶ Scheduler maintains wait-for-graph for detecting cycles
 - ▶ if transaction is blocked, add to graph
 - ▶ chain like T_i waits for T_j waits for T_k waits for T_i creates cycle

▶ **Deadlock avoidance**

- ▶ Avoid deadlock by not letting the cycle complete
 - ▶ When blocking transaction likely to create a cycle, abort the transaction

▶ **Starvation** means transaction consistently waits or restarts and never completes

- ▶ In deadlock resolution, same transaction may repeatedly be victim and aborted

Practical deadlock prevention



- ▶ Prevent deadlocks when a requested lock cannot be obtained
 - ▶ Based on the timestamp (“age”) of transaction: when did it enter?
 - ▶ **Wait-die**: older transaction may wait for younger transaction, younger transaction waiting for older transaction is aborted
 - ▶ Has not yet been running as long, expect fewer resources are wasted
 - ▶ **Wound-wait**: younger transaction may wait for older one, an older transaction will preempt the younger one, i.e., younger one aborted as well
 - ▶ Both techniques handle deadlocks and restart the younger transaction
 - ▶ Still, may abort transactions that are not actually deadlocked
 - ▶ **No waiting**: as soon as lock cannot be obtained, abort transaction, restart with some delay
 - ▶ No deadlocks, but may cause many unnecessary aborts and restarts
 - ▶ **Cautious waiting**: if lock cannot be obtained, abort only if the transaction that has the conflicting lock is waiting, else wait
 - ▶ Also, no deadlocks

Versions of 2PL

▶ **Basic**

- ▶ Transaction locks data items incrementally
 - ▶ may cause deadlock which is dealt with

▶ **Strict**

- ▶ Exclusive locks unlocked after terminating (commit / abort and rollback)
 - ▶ most commonly used two-phase locking algorithm

▶ **Rigorous**

- ▶ All unlocking (shared and exclusive) after terminating (commit / abort and rollback)

▶ **Conservative**

- ▶ Prevent deadlock by locking all desired data items before start of transaction
 - ▶ Often not realistic / efficient: as part of execution, determine which data items needed; if not requested from beginning, need to abort and start over



A working example: 2PL



2PL: growth phase (acquire locks),
then shrink phase (release locks)

T ₁	T ₂
read_lock (X); read_item (X);	
	read_lock (X); read_item (X);
read_lock(Y); write_lock (Y); unlock (X);	
	write_lock(X); write_item(X);
write_item (Y); unlock (Y);	
	read_lock(Y);
commit;	read_item (Y); write_lock (Y); unlock (X); write_item(Y); unlock (Y); commit;

2PL basic: unlock immediately
when entering 2nd phase

T ₁	T ₂
read_lock (X); read_item (X);	
	read_lock (X); read_item (X);
read_lock(Y); write_lock (Y); unlock (X);	
	write_lock(X); write_item(X);
write_item (Y); commit; unlock (Y);	
	read_lock(Y); read_item (Y); write_lock (Y); write_item(Y); commit; unlock (X); unlock (Y);

2PL strict: unlock exclusive locks
when terminating

T ₁	T ₂
read_lock (X); read_item (X);	
	read_lock (X); read_item (X);
read_lock(Y); write_lock (Y); write_item (Y); commit; unlock (X);	
	write_lock(X); write_item(X);
unlock (Y);	
	read_lock(Y); read_item (Y); write_lock (Y); write_item(Y); commit; unlock (X); unlock (Y);

2PL rigorous: unlock when
terminating

Aggressive versus Conservative Protocols

▶ **Aggressive Protocols**

- ▶ Proceed as quickly as possible even if situation may lead to aborts later
- ▶ Example: our practical 2-phase lock conversion protocol



▶ **Conservative protocols**

- ▶ Do not do work that may later have to be undone
- ▶ Example: each transaction requests all locks it will ever need at the beginning, and they are held until the end of the transactions
 - ▶ No deadlock!
 - ▶ Livelock (also called starvation) avoided by queuing lock requests
 - ▶ Disadvantage - more waiting, loss of concurrency



Intended learning outcomes

- ▶ Be able to
 - ▶ Determine serializable schedules
 - ▶ Apply the two-phase locking protocol

Acknowledgements: includes slide material by Jeff Ullman, Torben Bach Pedersen, and teachcoop

What was this all about?

Guidelines for your own review of today's session

- ▶ Concurrency control allows transactions to...
- ▶ A schedule is...
 - ▶ It is serial if...
- ▶ A conflict between transactions is...
 - ▶ A precedence graph shows conflicts by...
 - ▶ A (conflict-)serializable schedule is...
- ▶ The two-phase locking protocol consists of...
 - ▶ It guarantees that... because...
 - ▶ The difference between a shared and exclusive lock is...
- ▶ Deadlocks are...
 - ▶ They can happen if...
- ▶ Starvation means...
 - ▶ It occurs when...